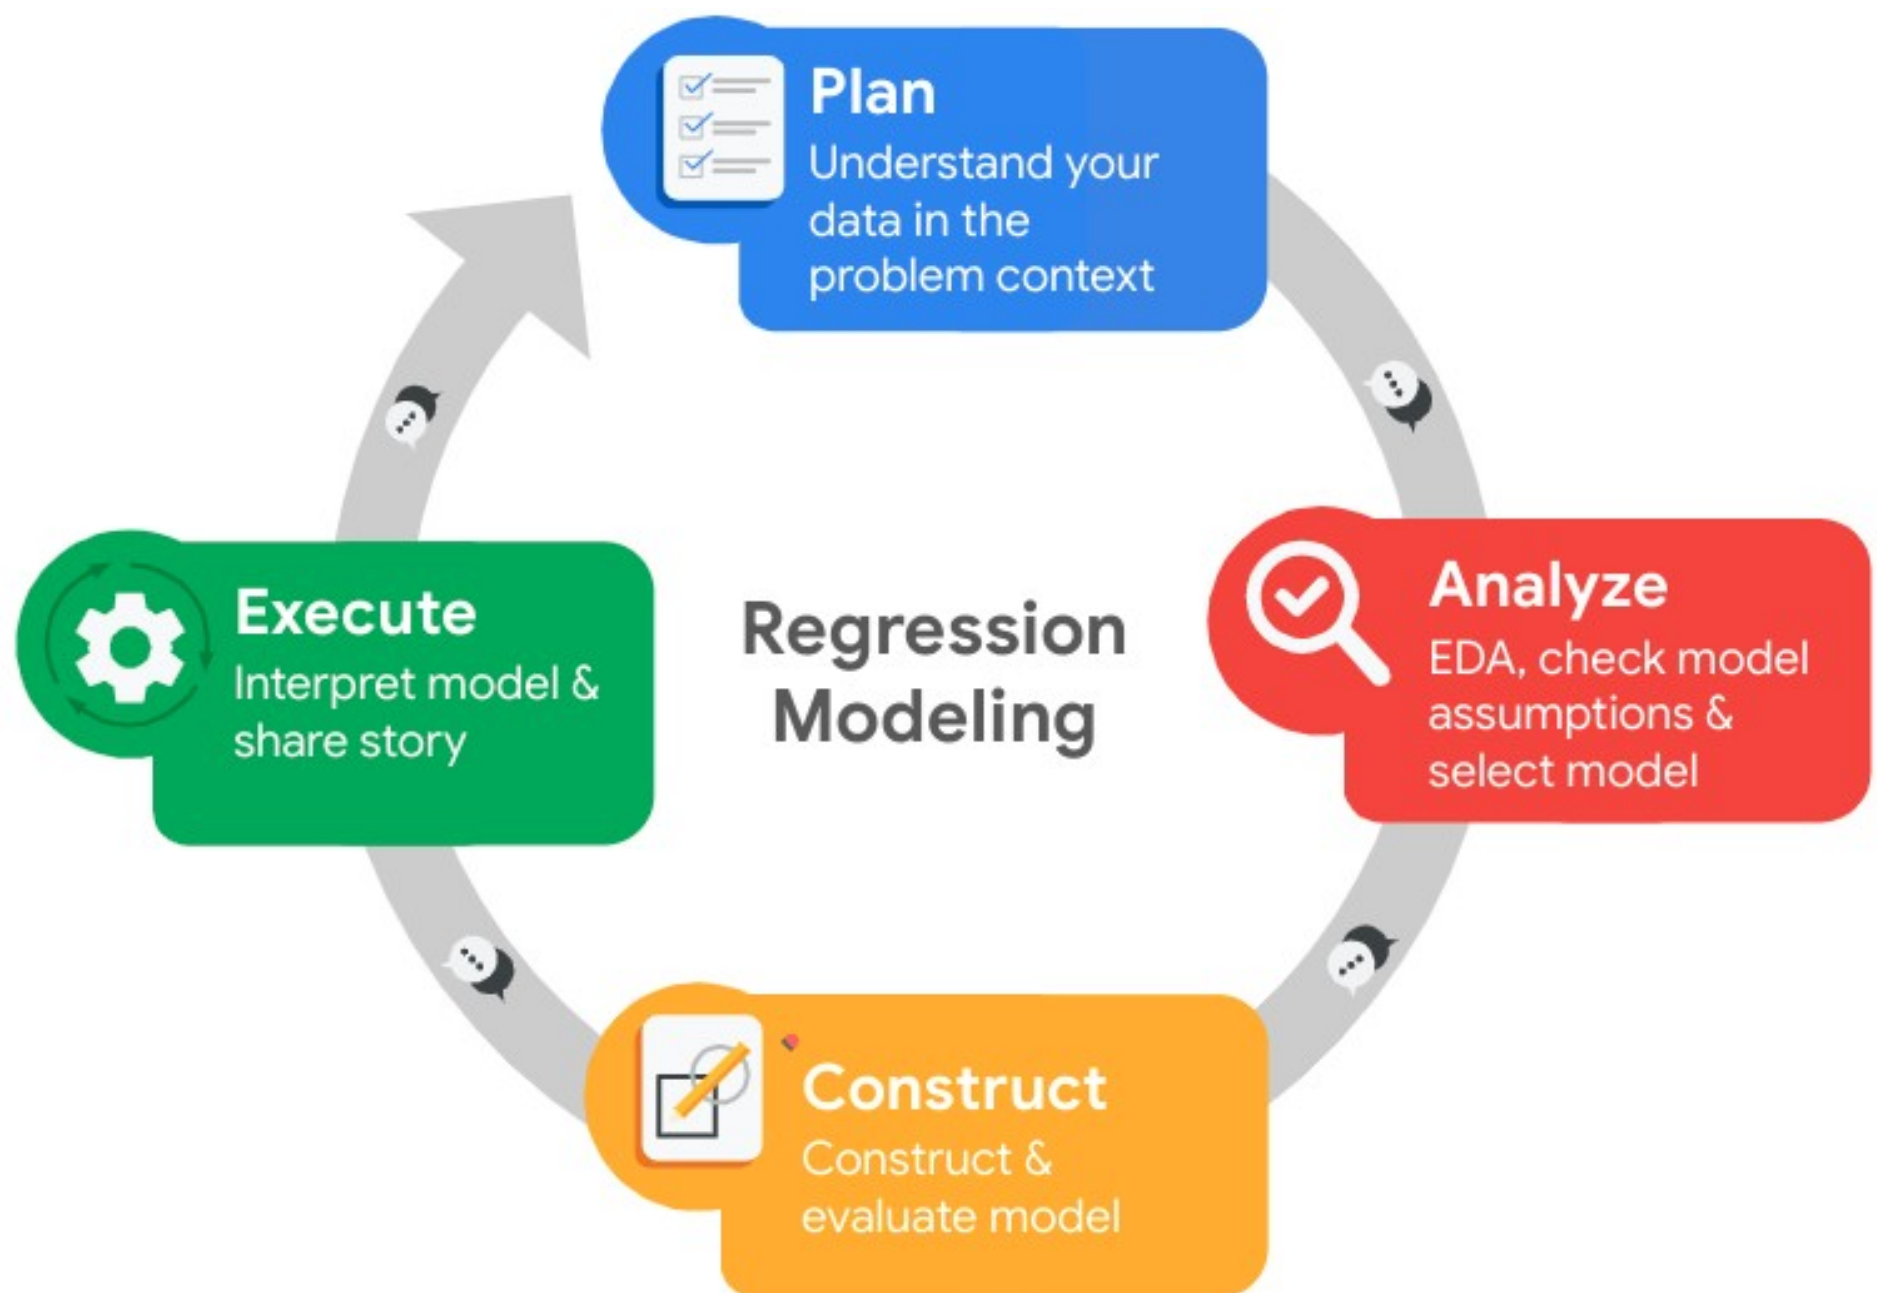


PACE Stages: The Project Framework



Pace: Plan Stage

Estimating Taxi Fares for Riders in advance:

1. Project goal:

In this fictional scenario, the New York City Taxi and Limousine Commission (TLC) has approached the data consulting firm Automatidata to develop an app that enables TLC riders to estimate the taxi fares in advance of their ride.

1. Background:

Since 1971, TLC has been regulating and overseeing the licensing of New York City's taxi cabs, for-hire vehicles, commuter vans, and paratransit vehicles.

1. Scenario:

New York City TLC stakeholders have been impressed with the data analytical work completed by the Automatidata team in this project. As a result, they have reached out once again for assistance in creating a machine learning model that can help predict the taxi fares for riders in advance.

Step 1. Imports & Loading Dataset

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from scipy import stats

# For displaying all of the columns in dataframes
pd.set_option('display.max_columns', None)

import warnings
warnings.filterwarnings('ignore')

from datetime import date, datetime, timedelta

from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
import sklearn.metrics as metrics

from xgboost import XGBRegressor, plot_importance
from sklearn.ensemble import RandomForestRegressor

df = pd.read_csv(r"D:\Google Advanced Data Analytics\Go Beyond the Numbers\Automatidata_EDA_Project\
2017_Yellow_Taxi_Trip_Data.csv")
df.head(5)
```

```
   Unnamed: 0  VendorID  tpep_pickup_datetime  tpep_dropoff_datetime \
0      24870114         2  03/25/2017 8:55:43 AM  03/25/2017 9:09:47 AM
1      35634249         1    04/11/2017 14:53    04/11/2017 15:19
2     106203690         1  12/15/2017 7:26:56 AM  12/15/2017 7:34:08 AM
3      38942136         2    05/07/2017 13:17    05/07/2017 13:48
4      30841670         2  04/15/2017 11:32:20 PM  04/15/2017 11:49:03 PM
```

```
   passenger_count  trip_distance  RatecodeID  store_and_fwd_flag \
0                6           3.34           1                N
1                1           1.80           1                N
2                1           1.00           1                N
3                1           3.70           1                N
4                1           4.37           1                N
```

```
   PULocationID  DOLocationID  payment_type  fare_amount  extra  mta_tax \
0            100           231            1          13.0     0.0    0.5
1            186            43            1          16.0     0.0    0.5
2            262           236            1           6.5     0.0    0.5
3            188            97            1          20.5     0.0    0.5
4             4           112            2          16.5     0.5    0.5
```

```
   tip_amount  tolls_amount  improvement_surcharge  total_amount
0         2.76           0.0                0.3          16.56
1         4.00           0.0                0.3          20.80
2         1.45           0.0                0.3           8.75
3         6.39           0.0                0.3          27.69
4         0.00           0.0                0.3          17.80
```

Step 2. Data Exploration (Initial EDA and data cleaning)

Gather basic information about the data

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            22699 non-null  int64
1   VendorID                              22699 non-null  int64
2   tpep_pickup_datetime                  22699 non-null  object
3   tpep_dropoff_datetime                 22699 non-null  object
4   passenger_count                       22699 non-null  int64
5   trip_distance                         22699 non-null  float64
6   RatecodeID                            22699 non-null  int64
7   store_and_fwd_flag                   22699 non-null  object
8   PULocationID                         22699 non-null  int64
9   DOLocationID                         22699 non-null  int64
10  payment_type                          22699 non-null  int64
11  fare_amount                           22699 non-null  float64
12  extra                                 22699 non-null  float64
13  mta_tax                               22699 non-null  float64
14  tip_amount                            22699 non-null  float64
15  tolls_amount                          22699 non-null  float64
16  improvement_surcharge                 22699 non-null  float64
17  total_amount                          22699 non-null  float64
dtypes: float64(8), int64(7), object(3)
memory usage: 3.1+ MB
```

Gather descriptive statistics about the data

```
df.describe(include= "all").T
```

```
   count  unique  top  freq \
Unnamed: 0  22699.0  NaN  NaN  NaN
VendorID    22699.0  NaN  NaN  NaN
tpep_pickup_datetime  22699  22435  02/01/2017 21:08  3
tpep_dropoff_datetime  22699  22451  11/08/2017 22:34  3
passenger_count  22699.0  NaN  NaN  NaN
trip_distance    22699.0  NaN  NaN  NaN
RatecodeID       22699.0  NaN  NaN  NaN
store_and_fwd_flag  22699  2  N  22600
PULocationID     22699.0  NaN  NaN  NaN
DOLocationID     22699.0  NaN  NaN  NaN
payment_type     22699.0  NaN  NaN  NaN
fare_amount      22699.0  NaN  NaN  NaN
extra            22699.0  NaN  NaN  NaN
```

mta_tax	22699.0	NaN	NaN	NaN
tip_amount	22699.0	NaN	NaN	NaN
tolls_amount	22699.0	NaN	NaN	NaN
improvement_surcharge	22699.0	NaN	NaN	NaN
total_amount	22699.0	NaN	NaN	NaN

	mean	std	min	25%	\
Unnamed: 0	56758486.171285	32744929.492148	12127.0	28520556.0	
VendorID	1.556236	0.496838	1.0	1.0	
tpep_pickup_datetime	NaN	NaN	NaN	NaN	
tpep_dropoff_datetime	NaN	NaN	NaN	NaN	
passenger_count	1.642319	1.285231	0.0	1.0	
trip_distance	2.913313	3.653171	0.0	0.99	
RatecodeID	1.043394	0.708391	1.0	1.0	
store_and_fwd_flag	NaN	NaN	NaN	NaN	
PULocationID	162.412353	66.633373	1.0	114.0	
DOLocationID	161.527997	70.139691	1.0	112.0	
payment_type	1.336887	0.496211	1.0	1.0	
fare_amount	13.026629	13.243791	-120.0	6.5	
extra	0.333275	0.463097	-1.0	0.0	
mta_tax	0.497445	0.039465	-0.5	0.5	
tip_amount	1.835781	2.800626	0.0	0.0	
tolls_amount	0.312542	1.399212	0.0	0.0	
improvement_surcharge	0.299551	0.015673	-0.3	0.3	
total_amount	16.310502	16.097295	-120.3	8.75	

	50%	75%	max
Unnamed: 0	56731504.0	85374524.0	113486300.0
VendorID	2.0	2.0	2.0
tpep_pickup_datetime	NaN	NaN	NaN
tpep_dropoff_datetime	NaN	NaN	NaN
passenger_count	1.0	2.0	6.0
trip_distance	1.61	3.06	33.96
RatecodeID	1.0	1.0	99.0
store_and_fwd_flag	NaN	NaN	NaN
PULocationID	162.0	233.0	265.0
DOLocationID	162.0	233.0	265.0
payment_type	1.0	2.0	4.0
fare_amount	9.5	14.5	999.99
extra	0.0	0.5	4.5
mta_tax	0.5	0.5	0.5
tip_amount	1.35	2.45	200.0
tolls_amount	0.0	0.0	19.1
improvement_surcharge	0.3	0.3	0.3
total_amount	11.8	17.8	1200.29

Check missing values

```
df.isna().sum()

Unnamed: 0      0
VendorID        0
tpep_pickup_datetime  0
tpep_dropoff_datetime  0
passenger_count    0
trip_distance      0
RatecodeID        0
store_and_fwd_flag  0
PULocationID      0
DOLocationID      0
payment_type      0
fare_amount       0
extra             0
mta_tax           0
tip_amount        0
tolls_amount      0
improvement_surcharge  0
total_amount      0
dtype: int64
```

Check duplicates

```
df.duplicated().sum()

0
```

pAce: Analyze Stage

Step 3. Data Exploration (Continue EDA)

Convert pickup & dropoff columns to datetime

```
df[["tpep_pickup_datetime", "tpep_dropoff_datetime"]].dtypes

tpep_pickup_datetime    object
tpep_dropoff_datetime    object
dtype: object
```

```
df["tpep_pickup_datetime"] = pd.to_datetime(df["tpep_pickup_datetime"], format= 'mixed')
df["tpep_dropoff_datetime"] = pd.to_datetime(df["tpep_dropoff_datetime"], format= 'mixed')
```

```
df[["tpep_pickup_datetime", "tpep_dropoff_datetime"]].dtypes
```

```
tpep_pickup_datetime    datetime64[ns]
tpep_dropoff_datetime   datetime64[ns]
dtype: object
```

Create duration column

```
df["trip_duration"] = (df["tpep_dropoff_datetime"] - df["tpep_pickup_datetime"]) / np.timedelta64(1, "m")
df["trip_duration"].describe()
```

```
count    22699.000000
mean      17.010830
std       61.996458
min      -17.000000
25%       6.750000
50%      11.000000
75%      18.258333
max      1440.000000
Name: trip_duration, dtype: float64
```

Check outliers

```
round(df.describe(include= [np.number], percentiles= [.5]).T, 1)
```

	count	mean	std	min	50%	\
Unnamed: 0	22699.0	56758486.2	32744929.5	12127.0	56731504.0	
VendorID	22699.0	1.6	0.5	1.0	2.0	
passenger_count	22699.0	1.6	1.3	0.0	1.0	
trip_distance	22699.0	2.9	3.7	0.0	1.6	
RatecodeID	22699.0	1.0	0.7	1.0	1.0	
PULocationID	22699.0	162.4	66.6	1.0	162.0	
DOLocationID	22699.0	161.5	70.1	1.0	162.0	
payment_type	22699.0	1.3	0.5	1.0	1.0	
fare_amount	22699.0	13.0	13.2	-120.0	9.5	
extra	22699.0	0.3	0.5	-1.0	0.0	
mta_tax	22699.0	0.5	0.0	-0.5	0.5	
tip_amount	22699.0	1.8	2.8	0.0	1.4	
tolls_amount	22699.0	0.3	1.4	0.0	0.0	
improvement_surcharge	22699.0	0.3	0.0	-0.3	0.3	
total_amount	22699.0	16.3	16.1	-120.3	11.8	
trip_duration	22699.0	17.0	62.0	-17.0	11.0	

	max
Unnamed: 0	113486300.0
VendorID	2.0
passenger_count	6.0
trip_distance	34.0
RatecodeID	99.0
PULocationID	265.0
DOLocationID	265.0
payment_type	4.0
fare_amount	1000.0
extra	4.5
mta_tax	0.5
tip_amount	200.0
tolls_amount	19.1
improvement_surcharge	0.3
total_amount	1200.3
trip_duration	1440.0

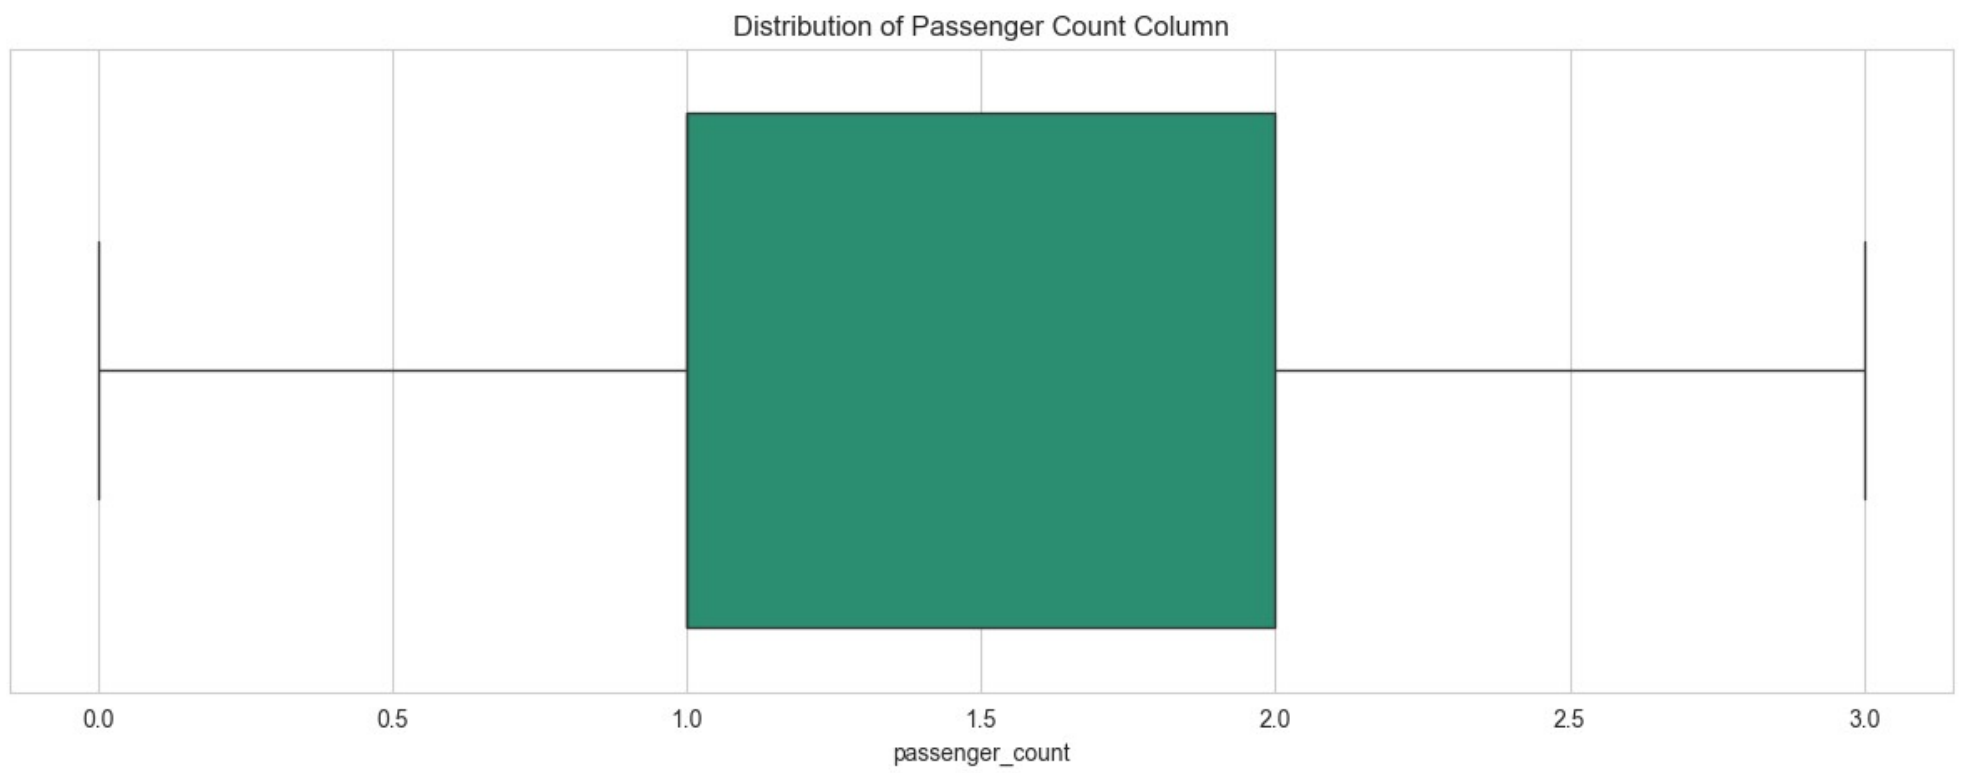
From above description for numerical variables, we could observe outliers in:

1. passenger_count
2. trip_distance
3. RatecodeID
4. fare_amount
5. extra
6. mta_tax
7. tip_amount
8. improvement_surcharge
9. total_amount
10. trip_duration

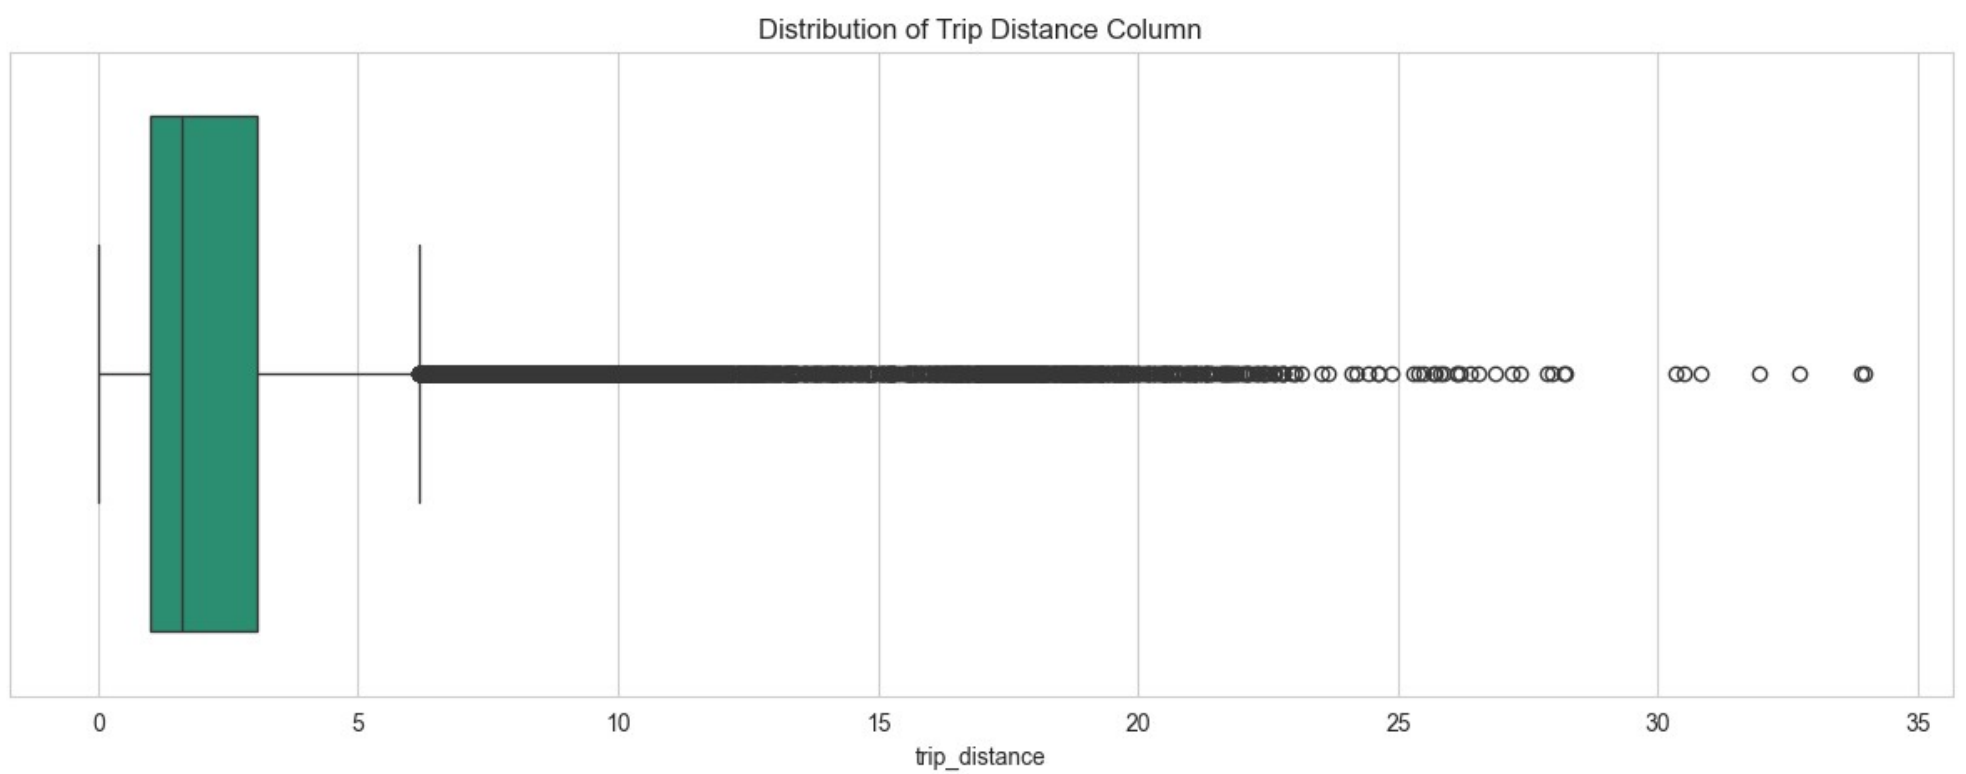
Box Plots

```
sns.set_style("whitegrid")
sns.set_palette("Dark2")
```

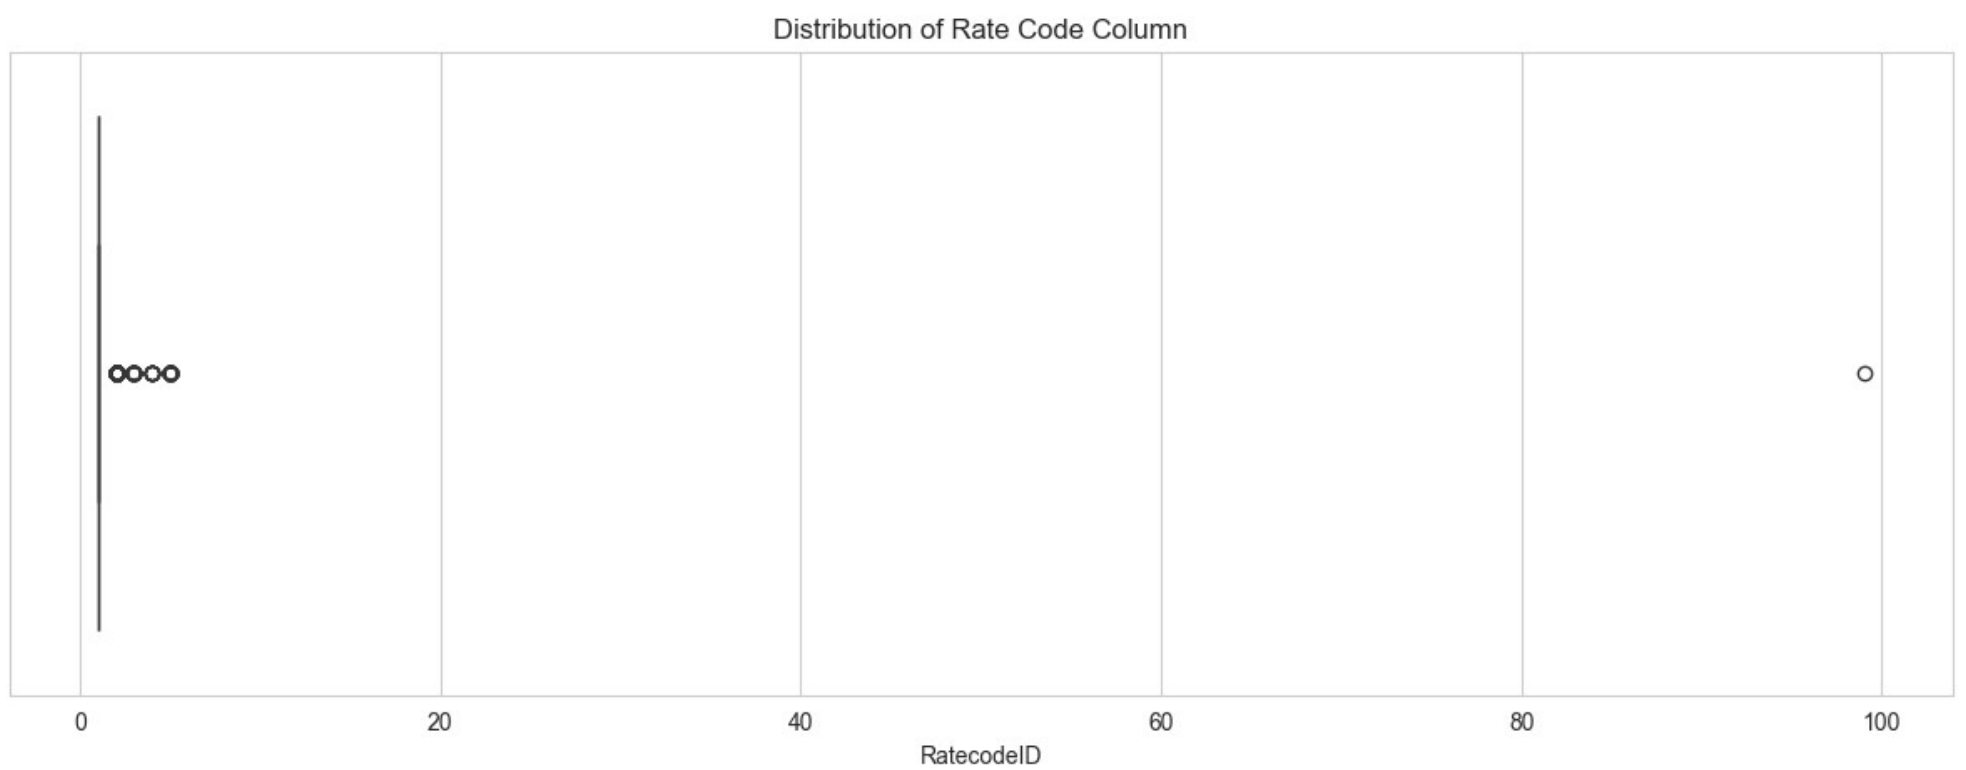
```
plt.figure(figsize = (15, 5))
sns.boxplot(data= df, x= "passenger_count", showfliers= False)
plt.title("Distribution of Passenger Count Column")
plt.show()
```



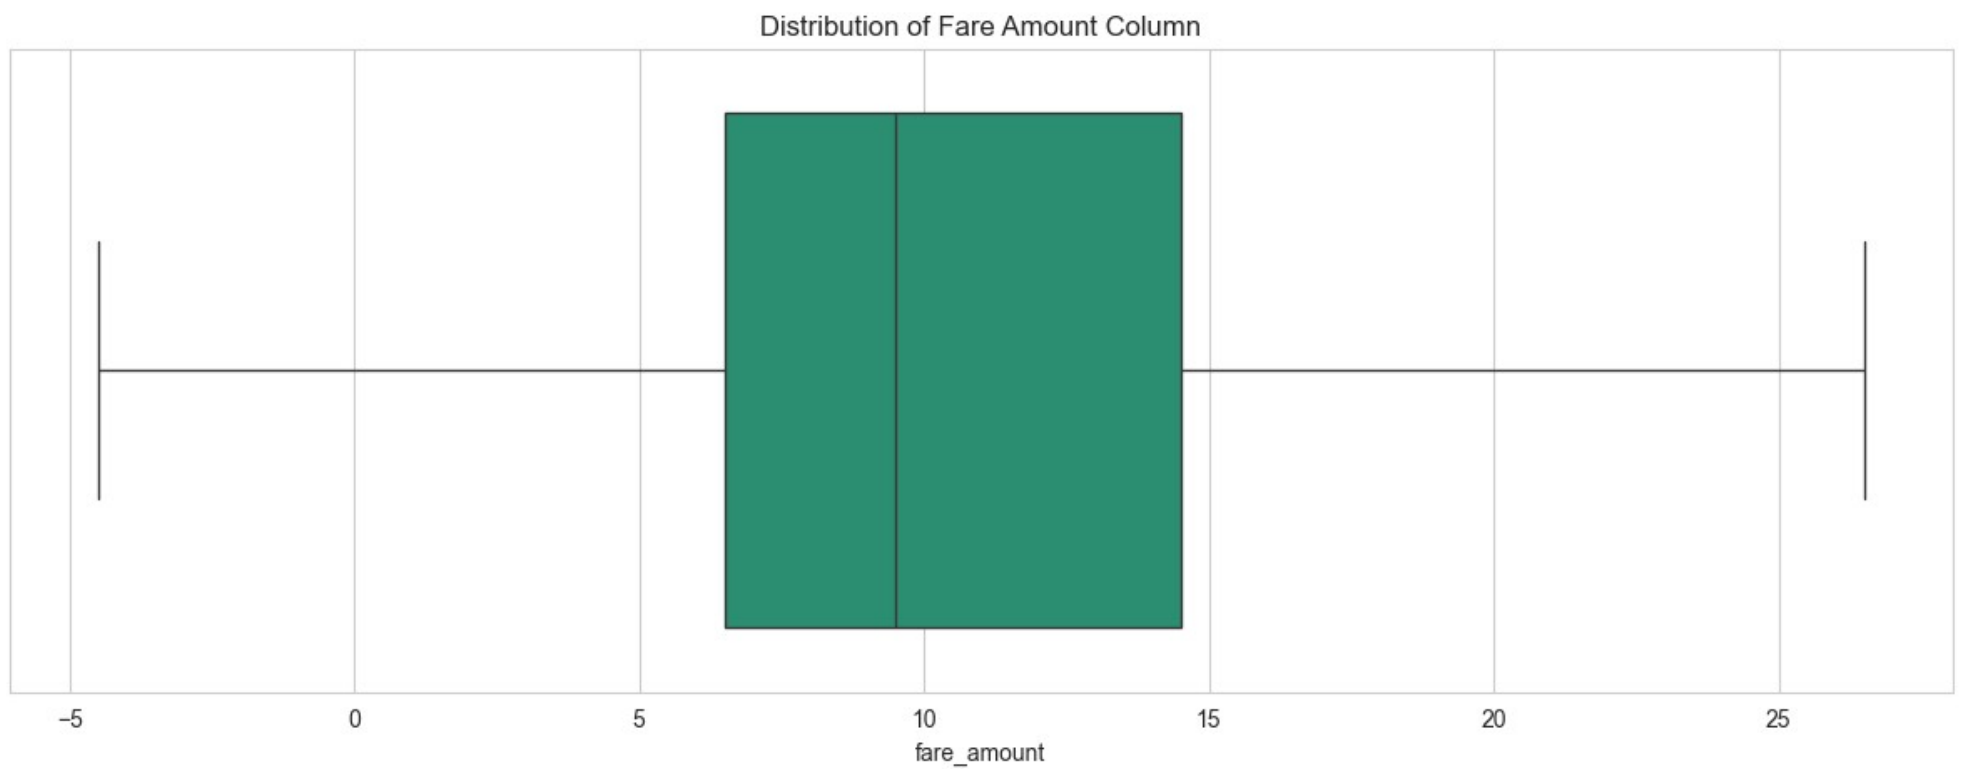
```
plt.figure(figsize = (15, 5))
sns.boxplot(data= df, x= "trip_distance", showliers= True)
plt.title("Distribution of Trip Distance Column")
plt.show()
```



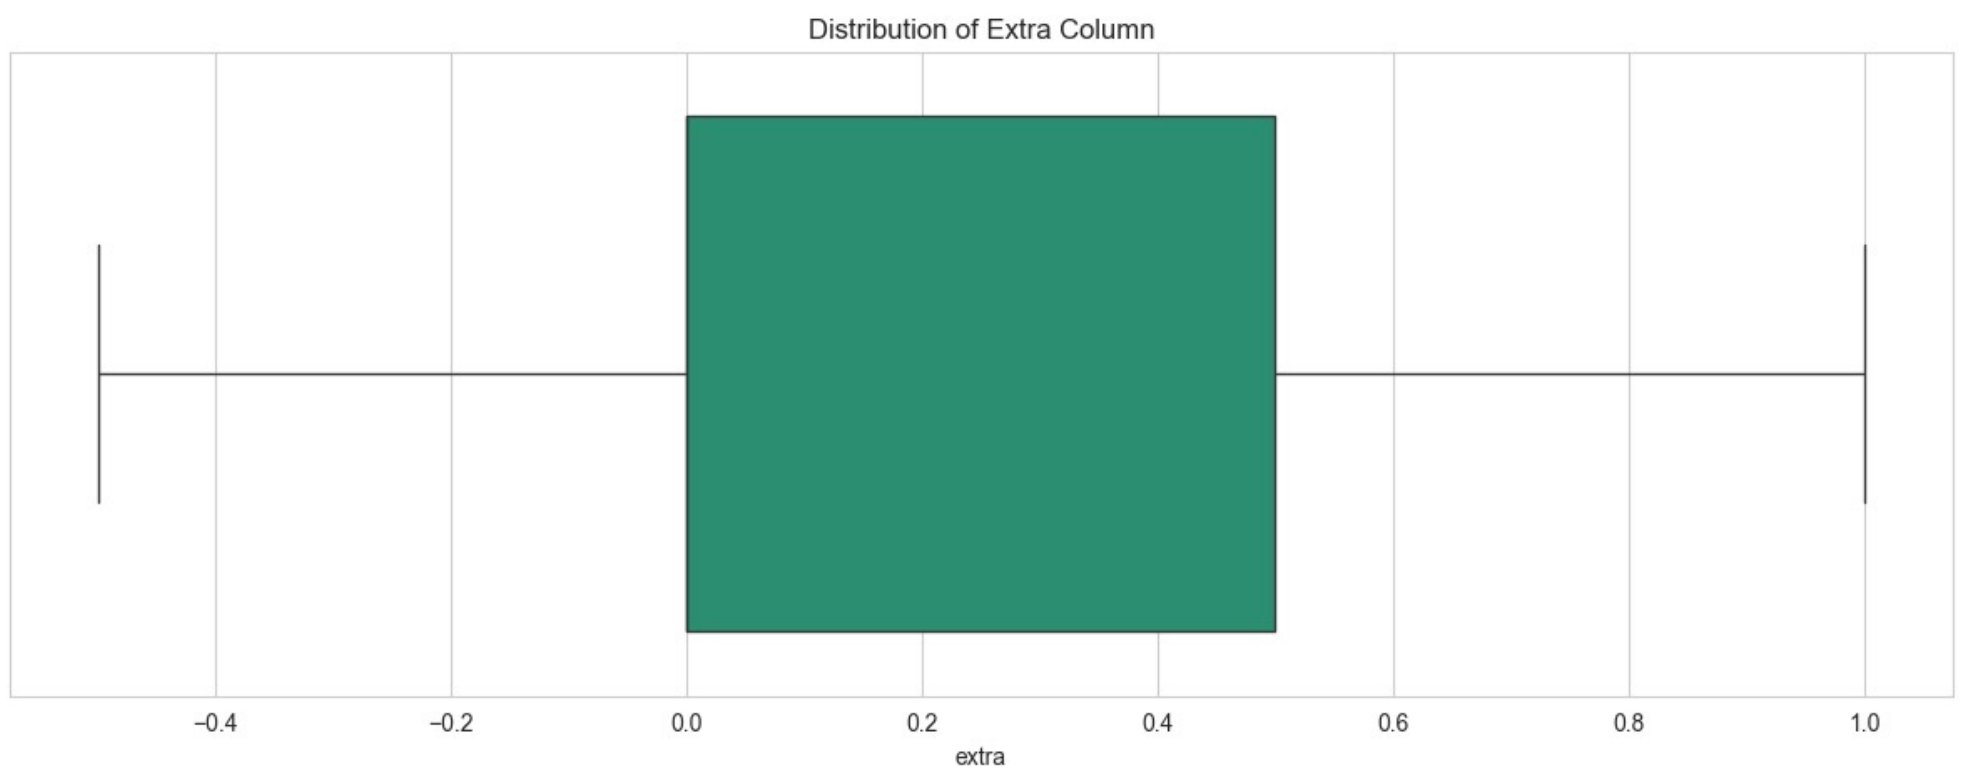
```
plt.figure(figsize = (15, 5))
sns.boxplot(data= df, x= "RatecodeID", showliers= True)
plt.title("Distribution of Rate Code Column")
plt.show()
```



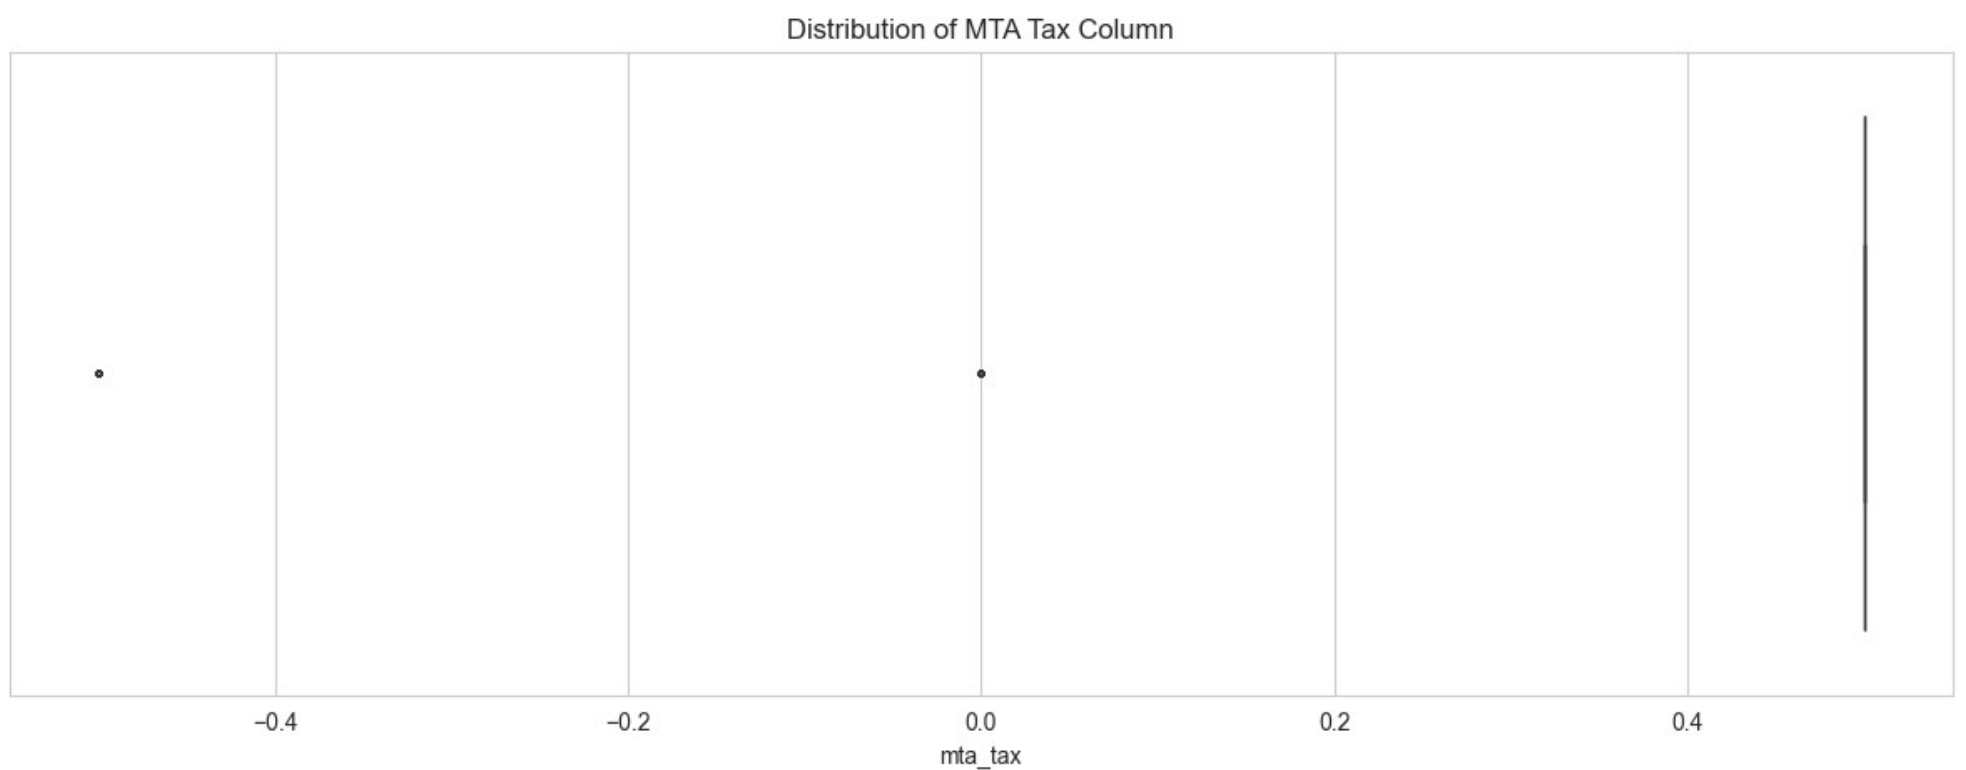
```
plt.figure(figsize = (15, 5))
sns.boxplot(data= df, x= "fare_amount", showliers= False)
plt.title("Distribution of Fare Amount Column")
plt.show()
```



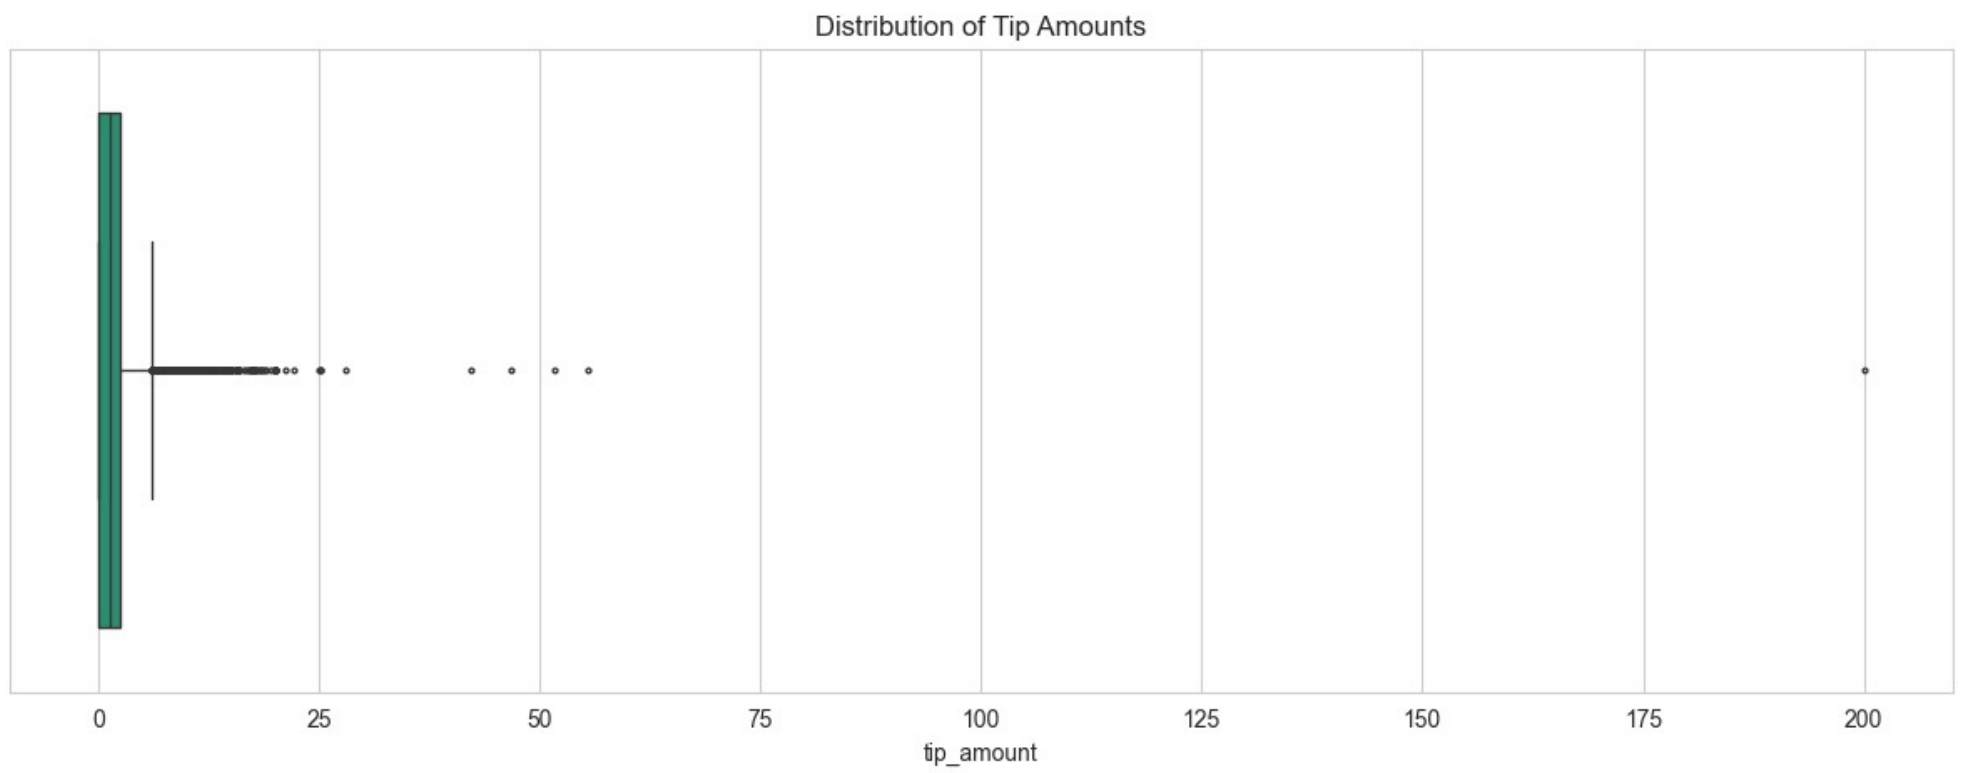
```
plt.figure(figsize = (15, 5))
sns.boxplot(data= df, x= "extra", showfliers= False)
plt.title("Distribution of Extra Column")
plt.show()
```



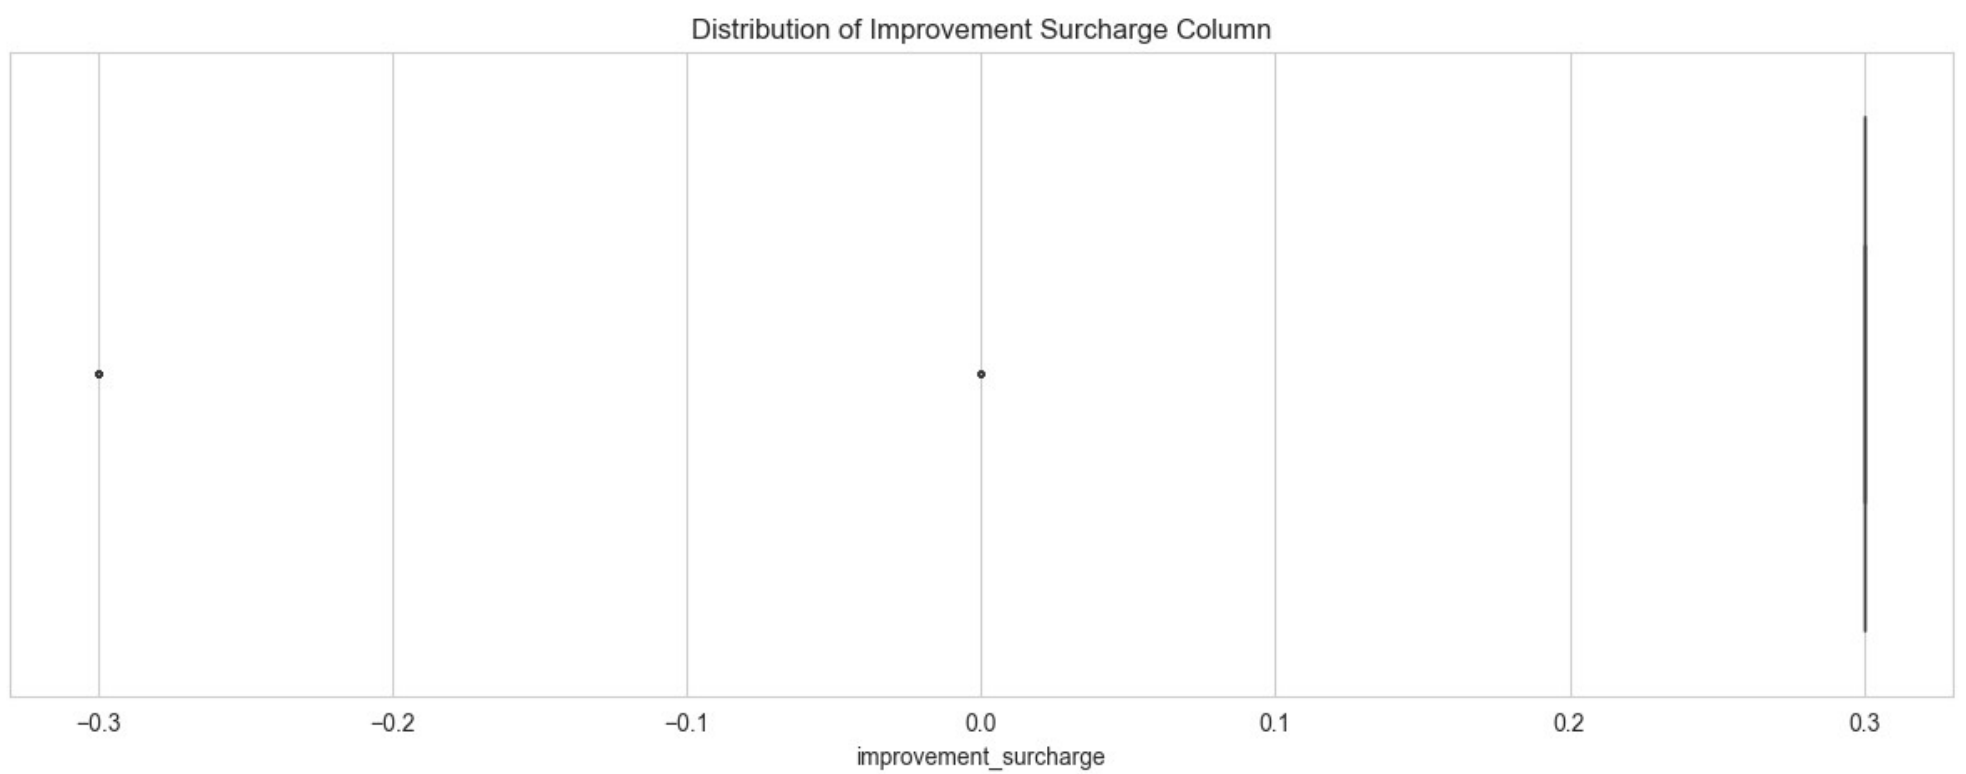
```
plt.figure(figsize = (15, 5))
sns.boxplot(data= df, x= "mta_tax", showfliers= True, fliersize= 2)
plt.title("Distribution of MTA Tax Column")
plt.show()
```




```
plt.figure(figsize = (15, 5))
sns.boxplot(data= df, x= "tip_amount", showfliers= True, fliersize= 2)
plt.title("Distribution of Tip Amounts")
plt.show()
```

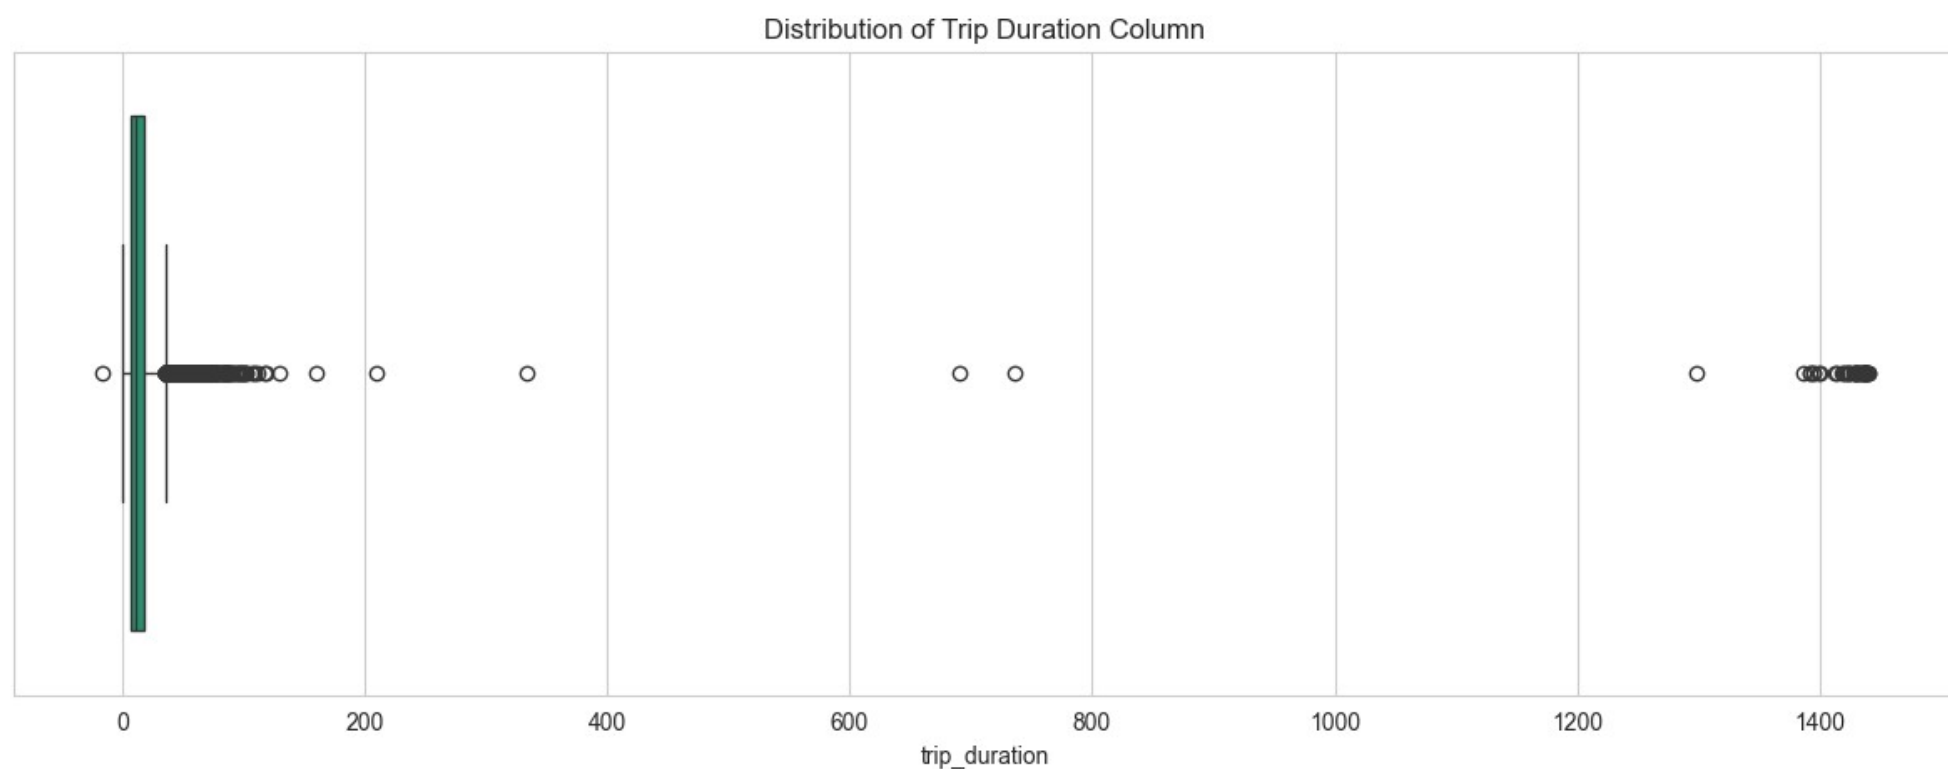
```
plt.figure(figsize = (15, 5))
sns.boxplot(data= df, x= "improvement_surcharge", showfliers= True, fliersize= 2)
plt.title("Distribution of Improvement Surcharge Column")
plt.show()
```



```
plt.figure(figsize = (15, 5))
sns.boxplot(data= df, x= "total_amount", showfliers= False)
plt.title("Distribution of Total Amount Column")
plt.show()
```



```
plt.figure(figsize = (15, 5))
sns.boxplot(data= df, x= "trip_duration", showfliers= True)
plt.title("Distribution of Trip Duration Column")
plt.show()
```



Imputations

passenger_count imputation

```
df[["passenger_count"]].value_counts()
```

passenger_count

1	16117
2	3305
5	1143
3	953
6	693
4	455
0	33

Name: count, dtype: int64

```
df.loc[df["passenger_count"] == 0, "passenger_count"] = 1
```

RatecodeID imputation

```
df[["RatecodeID"]].value_counts()
```

RatecodeID

1	22070
2	513
5	68
3	39
4	8
99	1

Name: count, dtype: int64

```
df.loc[df["RatecodeID"] == 99, "RatecodeID"] = 3
```

extra imputation

```
df.loc[df["extra"] < 0, "extra"].count()
```

9

```
df.loc[df["extra"] < 0, "extra"] = 0
```

mta_tax imputation

```
df.loc[df["mta_tax"] < 0, "mta_tax"].count()
```

13

```
df.loc[df["mta_tax"] < 0, "mta_tax"] = 0
```

tip_amount imputation

```
df["tip_amount"].sort_values(ascending= False).head(5)
```

8476	200.00
6064	55.50
13861	51.64
12511	46.69
984	42.29

Name: tip_amount, dtype: float64

```
df.loc[df["tip_amount"] == 200, "tip_amount"] = 55.50
```

improvement_surcharge imputation

```
df.loc[df["improvement_surcharge"] < 0, "improvement_surcharge"].count()
```


14

```
df.loc[df["improvement_surcharge"] < 0, "improvement_surcharge"] = 0
```

```
def impute_outliers(column_list, iqr_factor):  
    """  
    Impute upper-limit values in specified columns based on their interquartile range.
```

Arguments:

column_list: A list of columns to iterate over
iqr_factor: A number representing x in the formula:
 $Q3 + (x * IQR)$. Used to determine maximum threshold,
beyond which a point is considered an outlier.

The IQR is computed for each column in *column_list* and values exceeding the upper threshold for each column are imputed with the upper threshold value.

```
    """  
    for column in column_list:  
        # Reassign minimum to zero  
        df.loc[df[column] < 0, column] = 0  
  
        # Calculate upper threshold  
        Q3 = df[column].quantile(.75)  
        Q1 = df[column].quantile(.25)  
        iqr = Q3 - Q1  
        upper_limit = Q3 + (iqr_factor * iqr)  
        print(column)  
        print("Q3:", round(Q3, 3))  
        print("Upper threshold:", round(upper_limit, 3))  
  
        # Reassign values > threshold to threshold  
        df.loc[df[column] > upper_limit, column] = upper_limit
```

trip_distance imputation

```
impute_outliers(["trip_distance"], 9)
```

trip_distance

Q3: 3.06

Upper threshold: 21.69

```
df["trip_distance"].describe()
```

```
count    22699.000000  
mean      2.904191  
std       3.598010  
min       0.000000  
25%      0.990000  
50%      1.610000  
75%      3.060000  
max      21.690000  
Name: trip_distance, dtype: float64
```

trip_duration imputation

```
impute_outliers(["trip_duration"], 9)
```

trip_duration

Q3: 18.258

Upper threshold: 121.833

```
df["trip_duration"].describe()
```

```
count    22699.000000  
mean     14.538892  
std      12.546985  
min       0.000000  
25%      6.750000  
50%     11.000000  
75%     18.258333  
max     121.833333  
Name: trip_duration, dtype: float64
```

fare_amount imputation

```
impute_outliers(["fare_amount"], 9)
```

fare_amount

Q3: 14.5

Upper threshold: 86.5

```
df["fare_amount"].describe()
```

```
count    22699.000000  
mean     12.944139  
std      10.795620  
min       0.000000  
25%      6.500000
```

```
50%          9.500000
75%         14.500000
max          86.500000
Name: fare_amount, dtype: float64
```

total_amount imputation

```
impute_outliers(["total_amount"], 9)
```

```
total_amount
Q3: 17.8
Upper threshold: 99.25
```

```
df["total_amount"].describe()
```

```
count    22699.000000
mean      16.210212
std       13.337089
min        0.000000
25%        8.750000
50%       11.800000
75%       17.800000
max       99.250000
Name: total_amount, dtype: float64
```

Feature Engineering

Create 'pickup_dropoff' column

```
df["pickup_dropoff"] = df["PULocationID"].astype("str") + " >> " + df["DOLocationID"].astype("str")
df["pickup_dropoff"].head(5)
```

```
0    100 >> 231
1    186 >> 43
2    262 >> 236
3    188 >> 97
4      4 >> 112
Name: pickup_dropoff, dtype: object
```

```
df["pickup_dropoff"].describe()
```

```
count      22699
unique      4172
top    264 >> 264
freq        277
Name: pickup_dropoff, dtype: object
```

When deployed, the model will not know the duration of a trip until after the trip occurs, so you cannot train a model that uses this feature. However, you can use the statistics of trips you *do* know to generalize about ones you do not know.

Create mean_distance column

```
grouped_pickup_dropoff = round(df.groupby(["pickup_dropoff"])[["trip_distance"]].mean(), 2)
grouped_pickup_dropoff.head(5)
```

```
           trip_distance
pickup_dropoff
1 >> 1                2.43
10 >> 148             15.70
100 >> 1              16.89
100 >> 100             0.25
100 >> 107             1.18
```

1. Convert 'grouped' to a dictionary

```
grouped_dict = grouped_pickup_dropoff.to_dict()
```

2. Reassign to only contain the inner dictionary

```
grouped_dict = grouped_dict["trip_distance"]
```

1. Create a mean_distance column that is a copy of the pickup_dropoff helper column

```
df["mean_distance"] = df["pickup_dropoff"]
```

2. Map 'grouped_dict' to the 'mean_distance' column

```
df["mean_distance"] = df["mean_distance"].map(grouped_dict)
```

Confirm that it worked

```
df["mean_distance"].head(5)
```

```
0    3.52
1    3.11
2    0.88
3    3.70
4    4.44
Name: mean_distance, dtype: float64
```

Create mean_duration column

```
grouped_pickup_dropoff = round(df.groupby(["pickup_dropoff"])[["trip_duration"]].mean(), 2)
grouped_pickup_dropoff.head(5)
```

	trip_duration
pickup_dropoff	
1 >> 1	0.00
10 >> 148	69.00
100 >> 1	48.00
100 >> 100	3.12
100 >> 107	11.37

```
# 1. Convert `grouped` to a dictionary
```

```
grouped_dict = grouped_pickup_dropoff.to_dict()
```

```
# 2. Reassign to only contain the inner dictionary
```

```
grouped_dict = grouped_dict["trip_duration"]
```

```
# 1. Create a mean_duration column that is a copy of the pickup_dropoff helper column
```

```
df["mean_duration"] = df["pickup_dropoff"]
```

```
# 2. Map `grouped_dict` to the `mean_duration` column
```

```
df["mean_duration"] = df["mean_duration"].map(grouped_dict)
```

```
# Confirm that it worked
```

```
df["mean_duration"].head(5)
```

```
0    22.88
1    24.29
2     7.26
3    31.00
4    14.36
```

```
Name: mean_duration, dtype: float64
```

```
Create day & month columns
```

```
# Create 'day' col
```

```
df["day"] = df["tpep_pickup_datetime"].dt.day_name().str.lower()
```

```
# Create 'month' col
```

```
df["month"] = df["tpep_pickup_datetime"].dt.month_name().str.lower()
```

```
Create rush_hour column
```

Define rush hour as:

- Any weekday (not Saturday or Sunday) AND
- Either from 06:00–10:00 or from 16:00–20:00

Create a binary rush_hour column that contains a 1 if the ride was during rush hour and a 0 if it was not.

```
# Create 'rush_hour' col
```

```
df["rush_hour"] = df["tpep_pickup_datetime"].dt.hour
```

```
# If day is Saturday or Sunday, impute 0 in `rush_hour` column
```

```
df.loc[df["day"].isin(["saturday", "sunday"]), "rush_hour"] = 0
```

```
df["rush_hour"].head(5)
```

```
0    0
1   14
2    7
3    0
4    0
```

```
Name: rush_hour, dtype: int32
```

```
def rush_hourizer(hour):
```

```
    if 6 <= hour <= 10:
```

```
        val = 1
```

```
    elif 16 <= hour <= 20:
```

```
        val = 1
```

```
    else:
```

```
        val = 0
```

```
    return val
```

```
df["rush_hour"] = df["rush_hour"].apply(rush_hourizer)
```

```
df["rush_hour"].head(5)
```

```
0    0
1    0
2    1
3    0
4    0
```

```
Name: rush_hour, dtype: int64
```

```
Ordering columns
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 22699 entries, 0 to 22698
```

```
Data columns (total 25 columns):
```

```
#    Column                                Non-Null Count  Dtype
```

```

---
0  Unnamed: 0      22699 non-null int64
1  VendorID        22699 non-null int64
2  tpep_pickup_datetime 22699 non-null datetime64[ns]
3  tpep_dropoff_datetime 22699 non-null datetime64[ns]
4  passenger_count    22699 non-null int64
5  trip_distance      22699 non-null float64
6  RatecodeID        22699 non-null int64
7  store_and_fwd_flag 22699 non-null object
8  PULocationID       22699 non-null int64
9  DOLocationID       22699 non-null int64
10 payment_type       22699 non-null int64
11 fare_amount        22699 non-null float64
12 extra              22699 non-null float64
13 mta_tax            22699 non-null float64
14 tip_amount         22699 non-null float64
15 tolls_amount       22699 non-null float64
16 improvement_surcharge 22699 non-null float64
17 total_amount       22699 non-null float64
18 trip_duration      22699 non-null float64
19 pickup_dropoff     22699 non-null object
20 mean_distance      22699 non-null float64
21 mean_duration      22699 non-null float64
22 day                22699 non-null object
23 month              22699 non-null object
24 rush_hour          22699 non-null int64
dtypes: datetime64[ns](2), float64(11), int64(8), object(4)
memory usage: 4.3+ MB

```

```

df1 = df[["VendorID", "tpep_pickup_datetime", "tpep_dropoff_datetime", "day", "month", "rush_hour",\
         "passenger_count", "trip_distance", "trip_duration","RatecodeID", "store_and_fwd_flag",\
         "PULocationID", "DOLocationID", "pickup_dropoff", "mean_distance", "mean_duration", "payment_type",\
         "fare_amount", "extra", "mta_tax", "tip_amount", "tolls_amount", "improvement_surcharge",\
         "total_amount"]]
df1.shape

(22699, 24)

```

Data visualizations for tip_amount

Tips by Vendor

```
df1["VendorID"].value_counts(normalize= True)
```

```

VendorID
2    0.556236
1    0.443764
Name: proportion, dtype: float64

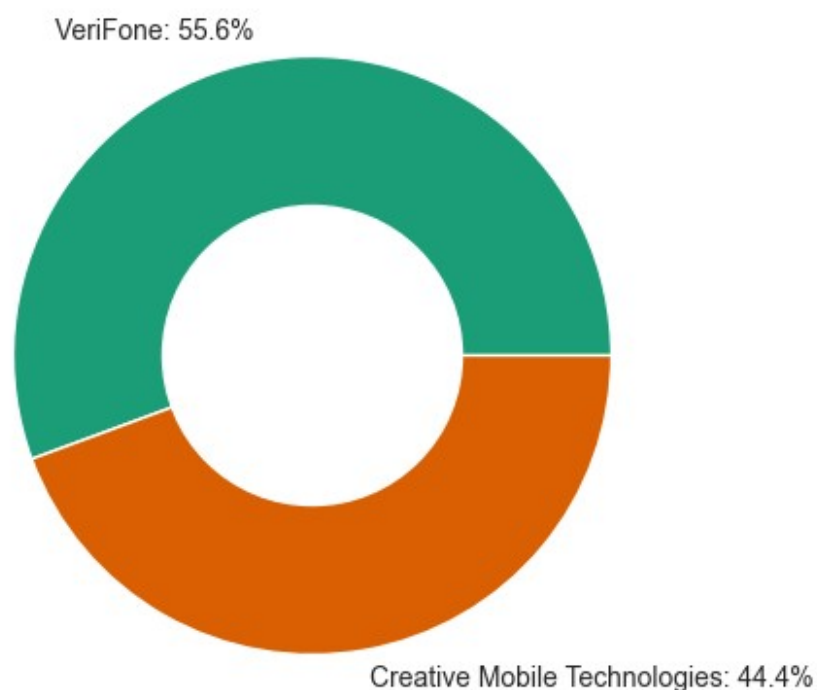
```

```

plt.figure(figsize= (15, 5))
plt.pie(df1["VendorID"].value_counts(), labels= ["VeriFone: 55.6%", "Creative Mobile Technologies: 44.4%"])

my_circle = plt.Circle( (0,0), 0.5, color= 'white')
p= plt.gcf()
p.gca().add_artist(my_circle)
plt.show()

```



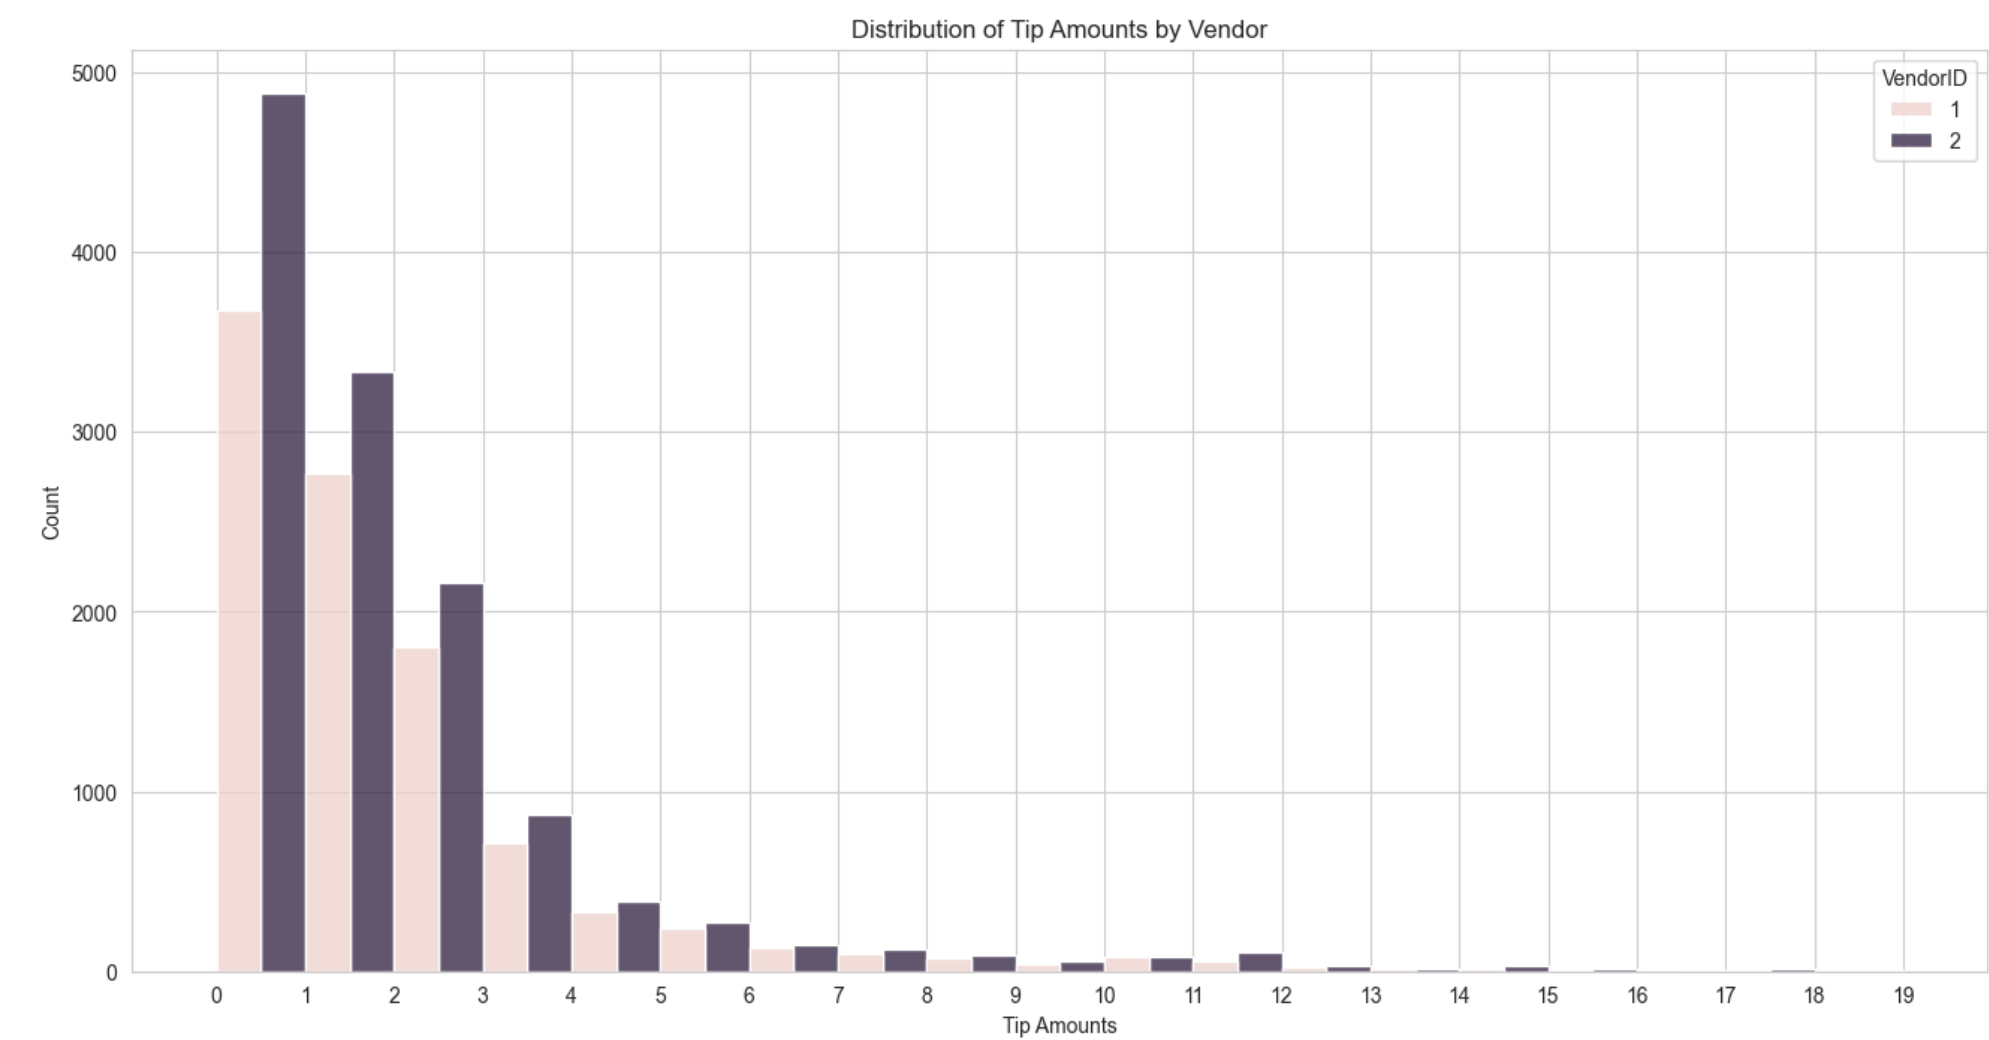
```

plt.figure(figsize = (16, 8))
ax= sns.histplot(data= df1, x= "tip_amount", bins= range(0, 20, 1), hue= "VendorID", multiple= "dodge")

ax.set_xticks(range(0, 20, 1))
ax.set_xticklabels(range(0, 20, 1))
plt.xlabel("Tip Amounts")

```

```
plt.title("Distribution of Tip Amounts by Vendor")
plt.show()
```



```
df_vendor = df1.groupby(["VendorID"])[["tip_amount", "total_amount"]].agg({"sum", "mean"})
df_vendor
```


VendorID	tip_amount		total_amount	
	sum	mean	sum	mean
1	18362.12	1.822905	162838.95	16.165884
2	23163.78	1.834610	205116.65	16.245577

```
Tips by Passenger Count
round(df1["passenger_count"].value_counts(normalize= True), 3) * 100
```

```
passenger_count
1    71.1
2    14.6
5     5.0
3     4.2
6     3.1
4     2.0
Name: proportion, dtype: float64
```

```
plt.figure(figsize= (15, 5))
plt.pie(df1["passenger_count"].value_counts(), labels= ["1: 71.1%", "2: 14.6%", "5: 5.0%", "3: 4.2%", "6: 3.1%", "4: 2.0%"])


my_circle = plt.Circle( (0,0), 0.5, color= 'white')
p= plt.gcf()
p.gca().add_artist(my_circle)
plt.show()
```



```
mean_tips_by_pass = df1.groupby(["passenger_count"])[["tip_amount"]].mean().reset_index().rename(columns=
{"tip_amount": "average_tip"})
mean_tips_by_pass
```

	passenger_count	average_tip
0	1	1.840559
1	2	1.856378
2	3	1.716768
3	4	1.530264
4	5	1.873185
5	6	1.720260

```
plt.figure(figsize = (15, 5))
sns.barplot(data= mean_tips_by_pass, x= "passenger_count", y= "average_tip", palette= "hls")
plt.title("Average of Tips by Passenger Count")
plt.show()
```



Tips by Payment Types


```
round(df1["payment_type"].value_counts(normalize= True), 3) * 100
```

```
payment_type
1    67.2
2    32.0
3     0.5
4     0.2
Name: proportion, dtype: float64
```

```
plt.figure(figsize= (15, 5))
plt.pie(df1["payment_type"].value_counts(), labels= ["Credit Card: 67.2%", "Cash: 32.0%", "Other: 0.8%", ""])

```

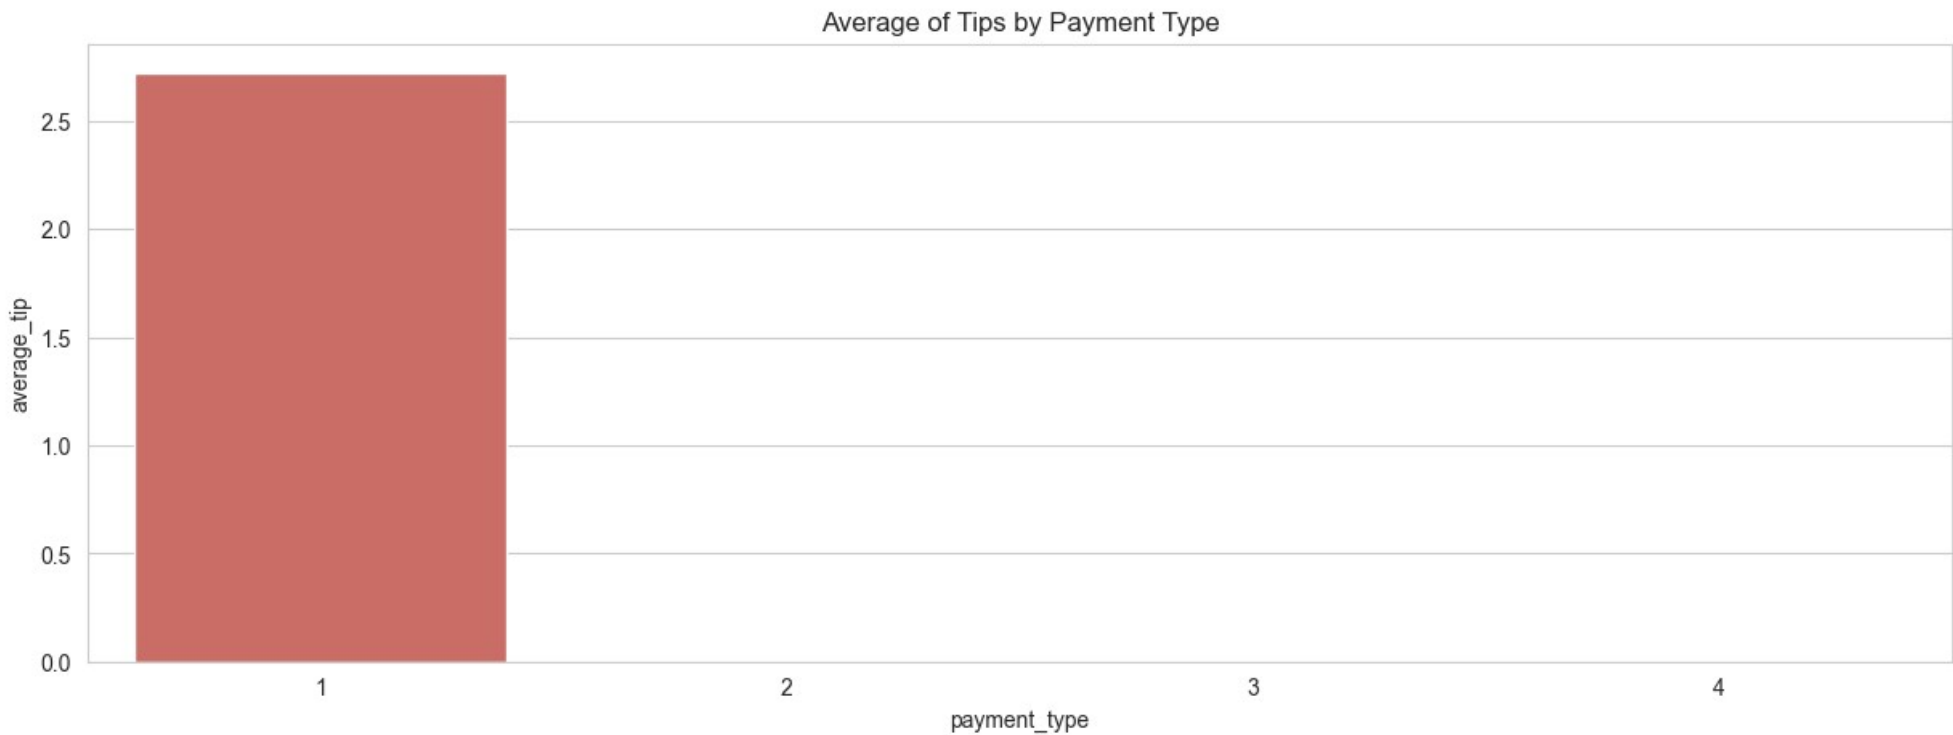
```
my_circle = plt.Circle( (0,0), 0.5, color= 'white')
p= plt.gcf()
p.gca().add_artist(my_circle)
plt.show()
```



```
mean_tips_by_pay = df1.groupby(["payment_type"])[["tip_amount"]].mean().reset_index().rename(columns={"tip_amount": "average_tip"})
mean_tips_by_pay

  payment_type  average_tip
0            1      2.720334
1            2      0.000000
2            3      0.000000
3            4      0.000000

plt.figure(figsize = (15, 5))
sns.barplot(data= mean_tips_by_pay, x= "payment_type", y= "average_tip", palette= "hls")
plt.title("Average of Tips by Payment Type")
plt.show()
```



Conduct an A/B Test with a two-sample t-test

```
df1.groupby(["payment_type"])[["fare_amount"]].agg({"mean", "median"})

      mean  median
payment_type
1    13.330766    9.5
2    12.141302    9.0
3    12.367934    7.0
4    12.989130    8.5
```

Based on the averages shown, it appears that customers who pay in credit card tend to pay a larger fare amount than customers who pay in cash. However, this difference might arise from rando sampling, rather than being a true difference in fare amount. To assess whether the difference s statistically significanwe'llyou conduct a hypothesis test.

```
credit_card = df1[df1["payment_type"] == 1]["fare_amount"]
cash = df1[df1["payment_type"] == 2]["fare_amount"]
print(len(credit_card))
print(len(cash))

15265
7267
```

Considering our hypotheses for this project as listed below:

- 1. H0: There is no difference in the average fare amount between customers who use credit cards and customers who use cash.

2. HA: There is a difference in the average fare amount between customers who use credit cards and customers who use cash.

We choose 5% as the significance level and proceed with a two-sample t-test.

```
stats.ttest_ind(a= credit_card, b= cash, equal_var= False)
```


```
TtestResult(statistic=7.928029528935512, pvalue=2.381707989138683e-15, df=15078.192451526904)
```

Since the p-value < significance level, then we reject the null hypothesis and accept the alternative hypothesis stating that there is a statistically significant difference in the average fare amount between customers who use credit cards and customers who use cash.

Scatter plots


```
plt.figure(figsize= (16, 8))
sns.scatterplot(data= df1, x= "mean_duration", y= "fare_amount")
plt.axhline(52, color= "red")
```

```
plt.xticks(range(0, 126, 5))
plt.yticks(range(0, 91, 5))
plt.title("Relationship between Mean Duration & Fare Amount")
plt.show()
```



```
plt.figure(figsize= (16, 8))
sns.scatterplot(data= df1, x= "mean_distance", y= "fare_amount")
plt.axhline(52, color= "red")
```

```
plt.xticks(range(0, 26, 5))
plt.yticks(range(0, 91, 5))
plt.title("Relationship between Mean Distance & Fare Amount")
plt.show()
```



The mean_duration variable correlates with the target variable. But what are the horizontal line around fare amount of 52 dollars?

```
df1[df1["fare_amount"] > 50]["fare_amount"].value_counts().head(3)
```

```
fare_amount
52.0      514
86.5       23
50.5        9
Name: count, dtype: int64
```

```
# Set pandas to display all columns
pd.set_option("display.max_columns", None)
df1[df1["fare_amount"] == 52].head(10)
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	day	\
11	2	2017-03-05 19:15:00	2017-03-05 19:52:00	sunday	
110	1	2017-06-03 14:24:00	2017-06-03 15:31:00	saturday	
161	2	2017-11-11 20:16:00	2017-11-11 20:17:00	saturday	
247	2	2017-12-06 23:37:00	2017-12-07 00:06:00	wednesday	
379	2	2017-09-24 23:45:45	2017-09-25 00:15:14	sunday	
388	1	2017-02-28 18:30:05	2017-02-28 19:09:55	tuesday	
406	2	2017-06-05 12:51:00	2017-06-05 13:07:00	monday	
449	2	2017-08-03 22:47:00	2017-08-03 23:32:00	thursday	
468	2	2017-09-26 13:48:26	2017-09-26 14:31:17	tuesday	
520	2	2017-04-23 21:34:48	2017-04-23 22:46:23	sunday	

	month	rush_hour	passenger_count	trip_distance	trip_duration	\
11	march	0	2	18.90	37.000000	
110	june	0	1	18.00	67.000000	
161	november	0	1	0.23	1.000000	
247	december	0	1	18.93	29.000000	
379	september	0	1	17.99	29.483333	
388	february	1	1	18.40	39.833333	
406	june	0	1	4.73	16.000000	
449	august	0	2	18.21	45.000000	
468	september	0	1	17.27	42.850000	
520	april	0	6	18.34	71.583333	

	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	pickup_dropoff	\
11	2	N	236	132	236 >> 132	
110	2	N	132	163	132 >> 163	
161	2	N	132	132	132 >> 132	
247	2	N	132	79	132 >> 79	
379	2	N	132	234	132 >> 234	
388	2	N	132	48	132 >> 48	
406	2	N	228	88	228 >> 88	
449	2	N	132	48	132 >> 48	
468	2	N	186	132	186 >> 132	
520	2	N	132	148	132 >> 148	

	mean_distance	mean_duration	payment_type	fare_amount	extra	mta_tax	\
11	19.21	46.10	1	52.0	0.0	0.5	
110	19.13	53.07	1	52.0	0.0	0.5	
161	2.16	3.02	2	52.0	0.0	0.5	

247	19.43	47.16	2	52.0	0.0	0.5
379	17.65	49.78	1	52.0	0.0	0.5
388	18.52	59.85	2	52.0	4.5	0.5
406	4.73	16.00	2	52.0	0.0	0.5
449	18.52	59.85	2	52.0	0.0	0.5
468	17.10	42.79	2	52.0	0.0	0.5
520	17.99	46.37	1	52.0	0.0	0.5

	tip_amount	tolls_amount	improvement_surcharge	total_amount
11	14.58	5.54	0.3	72.92
110	0.00	0.00	0.3	52.80
161	0.00	0.00	0.3	52.80
247	0.00	0.00	0.3	52.80
379	14.64	5.76	0.3	73.20
388	0.00	5.54	0.3	62.84
406	0.00	5.76	0.3	58.56
449	0.00	5.76	0.3	58.56
468	0.00	5.76	0.3	58.56
520	5.00	0.00	0.3	57.80

It seems that almost all of the trips in the first 30 rows where the fare amount was \$52 either begin or end at location 132, and all of them have a RatecodeID of 2.

There is no readily apparent reason why PULocation 132 should have so many fares of 52 dollars. They seem to occur on all different days, at different times, with both vendors, in all months. However, there are many toll amounts of \$5.76 and \ \$5.54. This would seem to indicate that location 132 is in an area that frequently requires tolls to get to and from. It's likely this is an airport.

Isolate modelling variables

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   VendorID                             22699 non-null  int64
1   tpep_pickup_datetime                 22699 non-null  datetime64[ns]
2   tpep_dropoff_datetime                 22699 non-null  datetime64[ns]
3   day                                   22699 non-null  object
4   month                                22699 non-null  object
5   rush_hour                             22699 non-null  int64
6   passenger_count                       22699 non-null  int64
7   trip_distance                         22699 non-null  float64
8   trip_duration                         22699 non-null  float64
9   RatecodeID                           22699 non-null  int64
10  store_and_fwd_flag                   22699 non-null  object
11  PULocationID                         22699 non-null  int64
12  DOLocationID                         22699 non-null  int64
13  pickup_dropoff                       22699 non-null  object
14  mean_distance                         22699 non-null  float64
15  mean_duration                         22699 non-null  float64
16  payment_type                          22699 non-null  int64
17  fare_amount                           22699 non-null  float64
18  extra                                 22699 non-null  float64
19  mta_tax                               22699 non-null  float64
20  tip_amount                           22699 non-null  float64
21  tolls_amount                          22699 non-null  float64
22  improvement_surcharge                 22699 non-null  float64
23  total_amount                          22699 non-null  float64
dtypes: datetime64[ns](2), float64(11), int64(7), object(4)
memory usage: 4.2+ MB
```

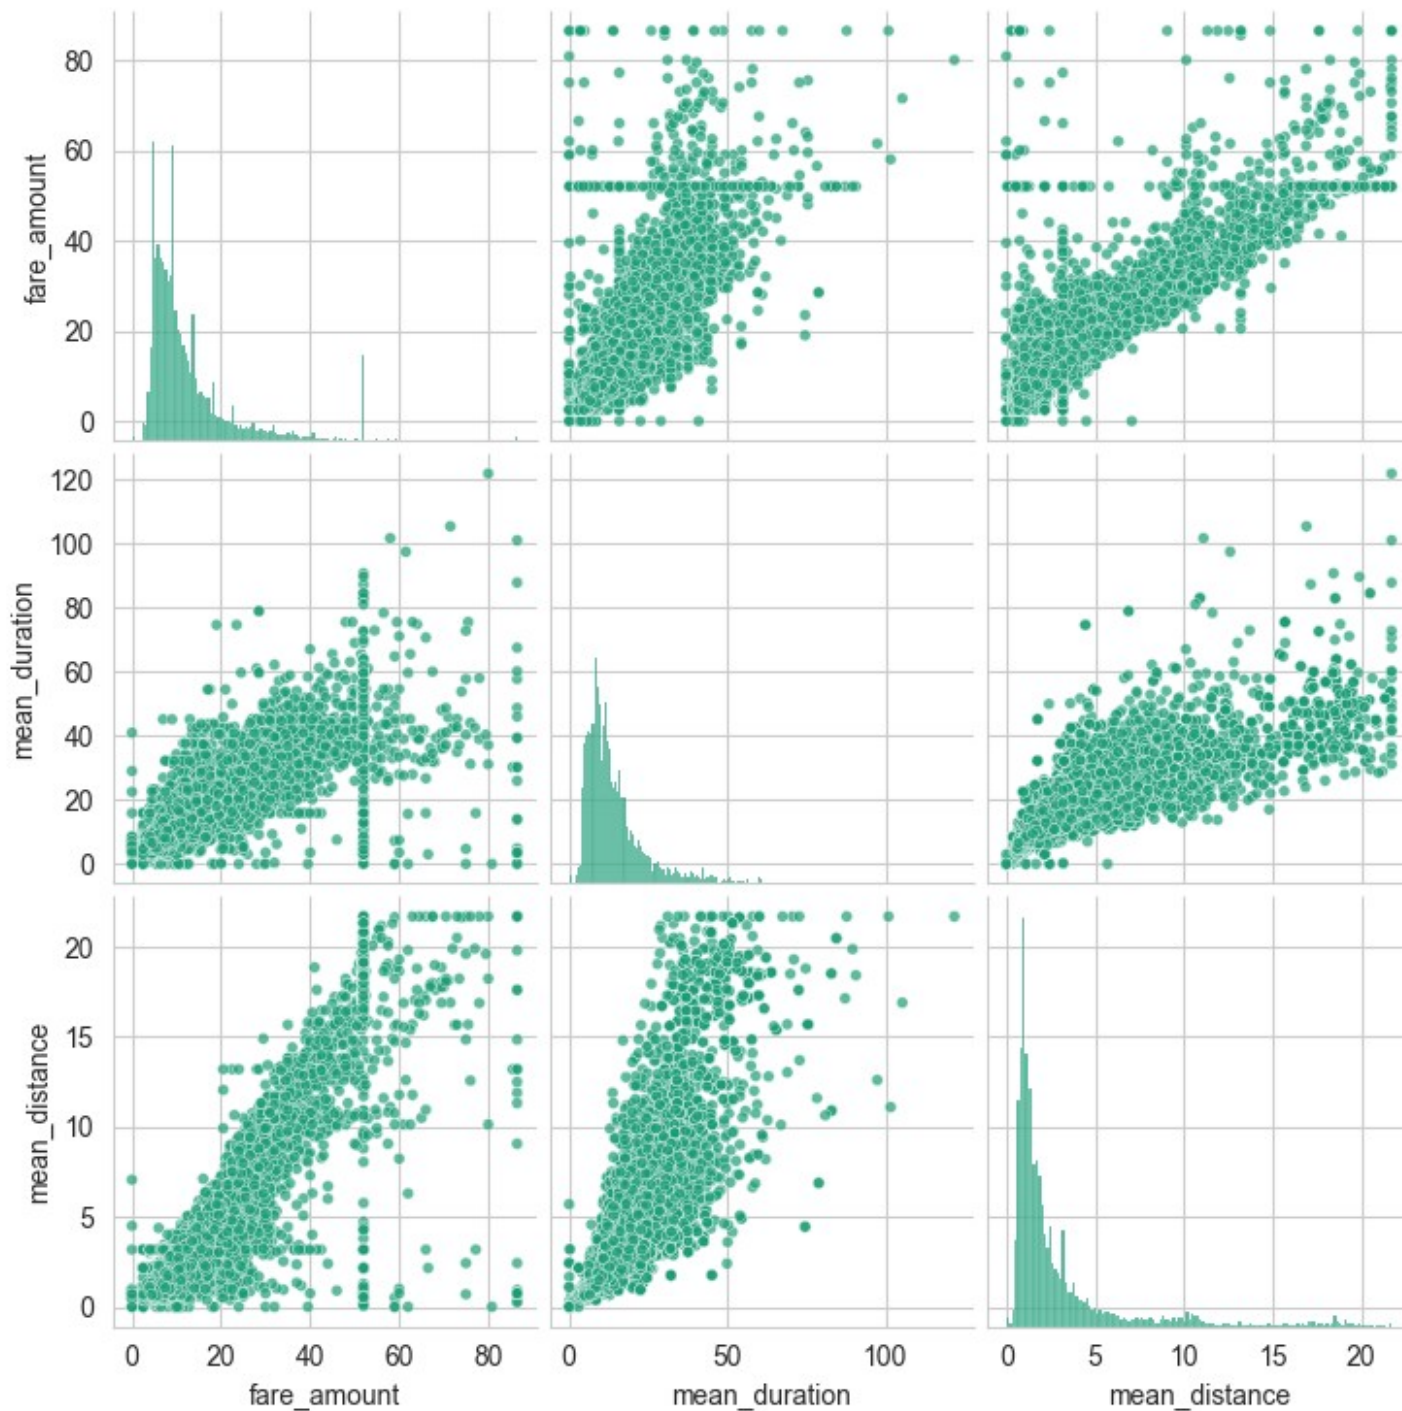
```
df2 = df1.copy()
```

```
df2 = df2.drop(["tpep_pickup_datetime", "tpep_dropoff_datetime", "day", "month", "trip_distance",
"trip_duration",\
                "store_and_fwd_flag", "PULocationID", "DOLocationID", "pickup_dropoff", "payment_type",
"RatecodeID",\
                "extra", "mta_tax", "tip_amount", "tolls_amount", "improvement_surcharge", "total_amount"],
axis= 1)
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   VendorID               22699 non-null  int64
1   rush_hour               22699 non-null  int64
2   passenger_count         22699 non-null  int64
3   mean_distance           22699 non-null  float64
4   mean_duration           22699 non-null  float64
5   fare_amount             22699 non-null  float64
dtypes: float64(3), int64(3)
memory usage: 1.0 MB
```


Pair plot

```
sns.pairplot(data= df2[["fare_amount", "mean_duration", "mean_distance"]], plot_kws= {"alpha": 0.7, "size": 0.7})  
plt.show()
```



Identify correlations


```
# Correlation matrix to help determine most correlated variables  
df2.corr(method= "pearson")
```

	VendorID	rush_hour	passenger_count	mean_distance	\
VendorID	1.000000	-0.000752	0.265464	0.003876	
rush_hour	-0.000752	1.000000	-0.024259	-0.047741	
passenger_count	0.265464	-0.024259	1.000000	0.013635	
mean_distance	0.003876	-0.047741	0.013635	1.000000	
mean_duration	0.002441	-0.026922	0.015732	0.869964	
fare_amount	0.001379	-0.025329	0.015012	0.902685	

	mean_duration	fare_amount
VendorID	0.002441	0.001379
rush_hour	-0.026922	-0.025329
passenger_count	0.015732	0.015012
mean_distance	0.869964	0.902685
mean_duration	1.000000	0.840942
fare_amount	0.840942	1.000000

Create correlation heatmap

```
plt.figure(figsize= (16, 8))  
sns.heatmap(data= df2.corr(method= "pearson"), vmin= -1, vmax= 1, annot= True, cmap= "YlGnBu")  
plt.show()
```



paCe: Construct Stage

Step 4. Model Building

Split data into outcome variable and features

```
# Remove the target column from the features
X = df2.drop(columns='fare_amount')
```

```
# Set y variable
```

```
y = df2[["fare_amount"]]
```

```
# Display first few rows
```

```
X.head(3)
```

	VendorID	rush_hour	passenger_count	mean_distance	mean_duration
0	2	0	6	3.52	22.88
1	1	0	1	3.11	24.29
2	1	1	1	0.88	7.26

Dummy encode categorical vendorID

```
# Convert VendorID to string
```

```
X["VendorID"] = X["VendorID"].astype("str")
```

```
# Get dummies
```

```
X = pd.get_dummies(X, drop_first=False)
```

```
X.head(3)
```

	rush_hour	passenger_count	mean_distance	mean_duration	VendorID_1	\
0	0	6	3.52	22.88	False	
1	0	1	3.11	24.29	True	
2	1	1	0.88	7.26	True	

```
VendorID_2
```

```
0 True
```

```
1 False
```

```
2 False
```

Split data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size= 0.2, random_state= 0)
```

Modelling Approach [A]: Multiple Linear Regression

Standardize the data

```
# Standardize the X variables
```

```
scaler = StandardScaler().fit(X_train)
```

```
X_train_scaled = scaler.transform(X_train)
```

```
print("X_train_scaled:\n", X_train_scaled)
```

```
X_train_scaled:
```

```
[[-0.77153979 -0.50468931  0.88394718  0.16609516 -0.89286563  0.89286563]
```

```
[ 1.29610943 -0.50468931 -0.60457634 -0.70337029 -0.89286563  0.89286563]
[-0.77153979  0.27243722 -0.48265917 -0.57192814  1.11998936 -1.11998936]
...
[-0.77153979 -0.50468931 -0.45430634 -0.66637176  1.11998936 -1.11998936]
[ 1.29610943 -0.50468931 -0.59607049 -0.84844347  1.11998936 -1.11998936]
[-0.77153979  1.82669027  0.84992379  1.11734684 -0.89286563  0.89286563]]
```

Fit the model

Fit your model to the training data

```
mlr = LinearRegression(n_jobs= -1)
mlr.fit(X_train_scaled, y_train)
```

```
LinearRegression(n_jobs=-1)
```

Evaluate the model

```
def evaluate_regressor(actual, predicted):
    print(f"R^2: {round(r2_score(actual, predicted), 2)}")
    print(f"MAE: {round(mean_absolute_error(actual, predicted), 2)}")
    print(f"MSE: {round(mean_squared_error(actual, predicted), 2)}")
    print(f"RMSE: {round(np.sqrt(mean_squared_error(actual, predicted)), 2)}")
```

```
y_pred_train_mlr = mlr.predict(X_train_scaled)
```

```
evaluate_regressor(y_train, y_pred_train_mlr)
```

```
R^2: 0.82
MAE: 2.25
MSE: 20.95
RMSE: 4.58
```

```
X_test_scaled = scaler.transform(X_test)
```

```
y_pred_test_mlr = mlr.predict(X_test_scaled)
```

```
evaluate_regressor(y_test, y_pred_test_mlr)
```

```
R^2: 0.85
MAE: 2.19
MSE: 16.56
RMSE: 4.07
```

Coefficients for multiple linear regression

```
print("Coefficients:\n", mlr.coef_)
```

```
Coefficients:
[[ 1.40972853e-01  2.84896083e-02  7.59366967e+00  2.46856907e+00
 -2.19545604e+13 -2.19545604e+13]]
```

```
coefficients = round(pd.DataFrame(data= mlr.coef_, columns= X.columns) * 100, 4)
coefficients
```

```
   rush_hour  passenger_count  mean_distance  mean_duration  VendorID_1 \
0    14.0973             2.849           759.367          246.8569 -2.195456e+15
```

```
   VendorID_2
0 -2.195456e+15
```

These coefficients represent the relative increase in fare_amount for every unit increased in standard deviation of these variables. To simplify the interpretation, let's calculate the standard deviation of the most effective variables:

```
print(X_train['mean_distance'].std())
```

```
print(7.5937 / X_train['mean_distance'].std())
```

```
3.52708203496952
2.15296948716012
```

Now you can make a more intuitive interpretation: for every 3.53 miles traveled, the fare increased by a mean of \\$7.59. Or, for every 1 mile traveled, the fare increased by a mean of \$2.15.

```
print(X_train['mean_duration'].std())
```

```
print(2.4686 / X_train['mean_duration'].std())
```

```
10.270961194893056
0.24034751501421708
```

For every 10.27 minutes traveled, the fare increased by a mean of \$2.47. Or, for every 1 minute traveled, the fare increased by a mean of \$0.24.

Modelling Approach [B]: Random Forest Regressor

```
ranfor_reg = RandomForestRegressor(random_state= 17)
```

```
cv_params = {"max_depth": [4, 6, 8, 10], "min_samples_leaf": [1, 2, 3, 4], "n_estimators": [100, 150, 200,
250]}
```

```
scoring = ["explained_variance", "neg_mean_absolute_error", "neg_mean_squared_error",
"neg_root_mean_squared_error"]
```

```
ranfor_cv = GridSearchCV(estimator= ranfor_reg, param_grid= cv_params, scoring= scoring, cv= 5, refit=
"explained_variance", n_jobs= -1, verbose= 2)
```

```
%%time
```

```
ranfor_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 64 candidates, totalling 320 fits

CPU times: total: 4.48 s

Wall time: 3min 7s

```
GridSearchCV(cv=5, estimator=RandomForestRegressor(random_state=17), n_jobs=-1,
             param_grid={'max_depth': [4, 6, 8, 10],
                          'min_samples_leaf': [1, 2, 3, 4],
                          'n_estimators': [100, 150, 200, 250]},
             refit='explained_variance',
             scoring=['explained_variance', 'neg_mean_absolute_error',
                     'neg_mean_squared_error', 'neg_root_mean_squared_error'],
             verbose=2)
```

```
ranfor_cv.best_params_
```

```
{'max_depth': 8, 'min_samples_leaf': 1, 'n_estimators': 200}
```

```
ranfor_cv.best_score_
```

```
0.8391621789484782
```

Fit the model

```
ranfor_reg = RandomForestRegressor(n_estimators= 200, criterion= 'squared_error', max_depth= 8,
min_samples_leaf= 1, n_jobs= -1, verbose= 1, random_state= 17)
ranfor_reg.fit(X_train, y_train)
```

```
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:    0.0s
[Parallel(n_jobs=-1)]: Done 184 tasks    | elapsed:    0.5s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:    0.6s finished
```

```
RandomForestRegressor(max_depth=8, n_estimators=200, n_jobs=-1, random_state=17,
                      verbose=1)
```

Evaluate the model

Evaluate the model performance on the training data

```
y_pred_train_ranfor = ranfor_reg.predict(X_train)
```

```
evaluate_regressor(y_train, y_pred_train_ranfor)
```

```
R^2: 0.88
MAE: 2.05
MSE: 14.45
RMSE: 3.8
```

```
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:    0.0s
[Parallel(n_jobs=8)]: Done 184 tasks    | elapsed:    0.0s
[Parallel(n_jobs=8)]: Done 200 out of 200 | elapsed:    0.0s finished
```

Evaluate the model performance on the testing data

```
y_pred_test_ranfor = ranfor_reg.predict(X_test)
```

```
evaluate_regressor(y_test, y_pred_test_ranfor)
```

```
R^2: 0.86
MAE: 2.13
MSE: 15.41
RMSE: 3.93
```

```
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:    0.0s
[Parallel(n_jobs=8)]: Done 184 tasks    | elapsed:    0.0s
[Parallel(n_jobs=8)]: Done 200 out of 200 | elapsed:    0.0s finished
```

Modelling Approach [C]: Extreme Gradient Boost Regressor

```
xgboost_reg = XGBRegressor(objective= 'reg:squarederror', random_state= 17)
```

```
cv_params = {"max_depth": [4, 6, 8, 10], "min_child_weight": [1, 2, 3, 4], "n_estimators": [100, 150, 200,
250], "learning_rate": [0.05, 0.1, 0.2]}
```

```
scoring = ["explained_variance", "neg_mean_absolute_error", "neg_mean_squared_error",
"neg_root_mean_squared_error"]
```

```
xgboost_cv = GridSearchCV(estimator= xgboost_reg, param_grid= cv_params, scoring= scoring, cv= 5, refit=
"explained_variance", n_jobs= -1, verbose= 2)
```



```
%%time
```

```
xgboost_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 192 candidates, totalling 960 fits

CPU times: total: 8.98 s

Wall time: 1min 31s

```
GridSearchCV(cv=5,
             estimator=XGBRegressor(base_score=None, booster=None,
                                     callbacks=None, colsample_bylevel=None,
                                     colsample_bynode=None,
                                     colsample_bytree=None, device=None,
                                     early_stopping_rounds=None,
                                     enable_categorical=False, eval_metric=None,
                                     feature_types=None, gamma=None,
                                     grow_policy=None, importance_type=None,
                                     interaction_constraints=None,
                                     learning_rate=None, m...
                                     multi_strategy=None, n_estimators=None,
                                     n_jobs=None, num_parallel_tree=None,
                                     random_state=17, ...),
             n_jobs=-1,
             param_grid={'learning_rate': [0.05, 0.1, 0.2],
                         'max_depth': [4, 6, 8, 10],
                         'min_child_weight': [1, 2, 3, 4],
                         'n_estimators': [100, 150, 200, 250]},
             refit='explained_variance',
             scoring=['explained_variance', 'neg_mean_absolute_error',
                     'neg_mean_squared_error', 'neg_root_mean_squared_error'],
             verbose=2)
```

```
xgboost_cv.best_params_
```

```
{'learning_rate': 0.05,
 'max_depth': 4,
 'min_child_weight': 3,
 'n_estimators': 250}
```

```
xgboost_cv.best_score_
```

```
0.8364500917599068
```

Fit the model

```
xgboost_reg = XGBRegressor(n_estimators= 250, learning_rate= 0.05, max_depth= 4, min_child_weight= 3, n_jobs= -
1, verbose= 1, random_state= 17)
xgboost_reg.fit(X_train, y_train)
```

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, device=None, early_stopping_rounds=None,
             enable_categorical=False, eval_metric=None, feature_types=None,
             gamma=None, grow_policy=None, importance_type=None,
             interaction_constraints=None, learning_rate=0.05, max_bin=None,
             max_cat_threshold=None, max_cat_to_onehot=None,
             max_delta_step=None, max_depth=4, max_leaves=None,
             min_child_weight=3, missing=nan, monotone_constraints=None,
             multi_strategy=None, n_estimators=250, n_jobs=-1,
             num_parallel_tree=None, random_state=17, ...)
```

Evaluate the model

```
# Evaluate the model performance on the training data
```

```
y_pred_train_xgb = xgboost_reg.predict(X_train)
```

```
evaluate_regressor(y_train, y_pred_train_xgb)
```

```
R^2: 0.86
```

```
MAE: 2.12
```

```
MSE: 16.2
```

```
RMSE: 4.03
```

```
# Evaluate the model performance on the testing data
```

```
y_pred_test_xgb = xgboost_reg.predict(X_test)
```

```
evaluate_regressor(y_test, y_pred_test_xgb)
```

```
R^2: 0.86
```

```
MAE: 2.12
```

```
MSE: 15.55
```

```
RMSE: 3.94
```

pacE: Execute Stage

Step 5. Results and Evaluation

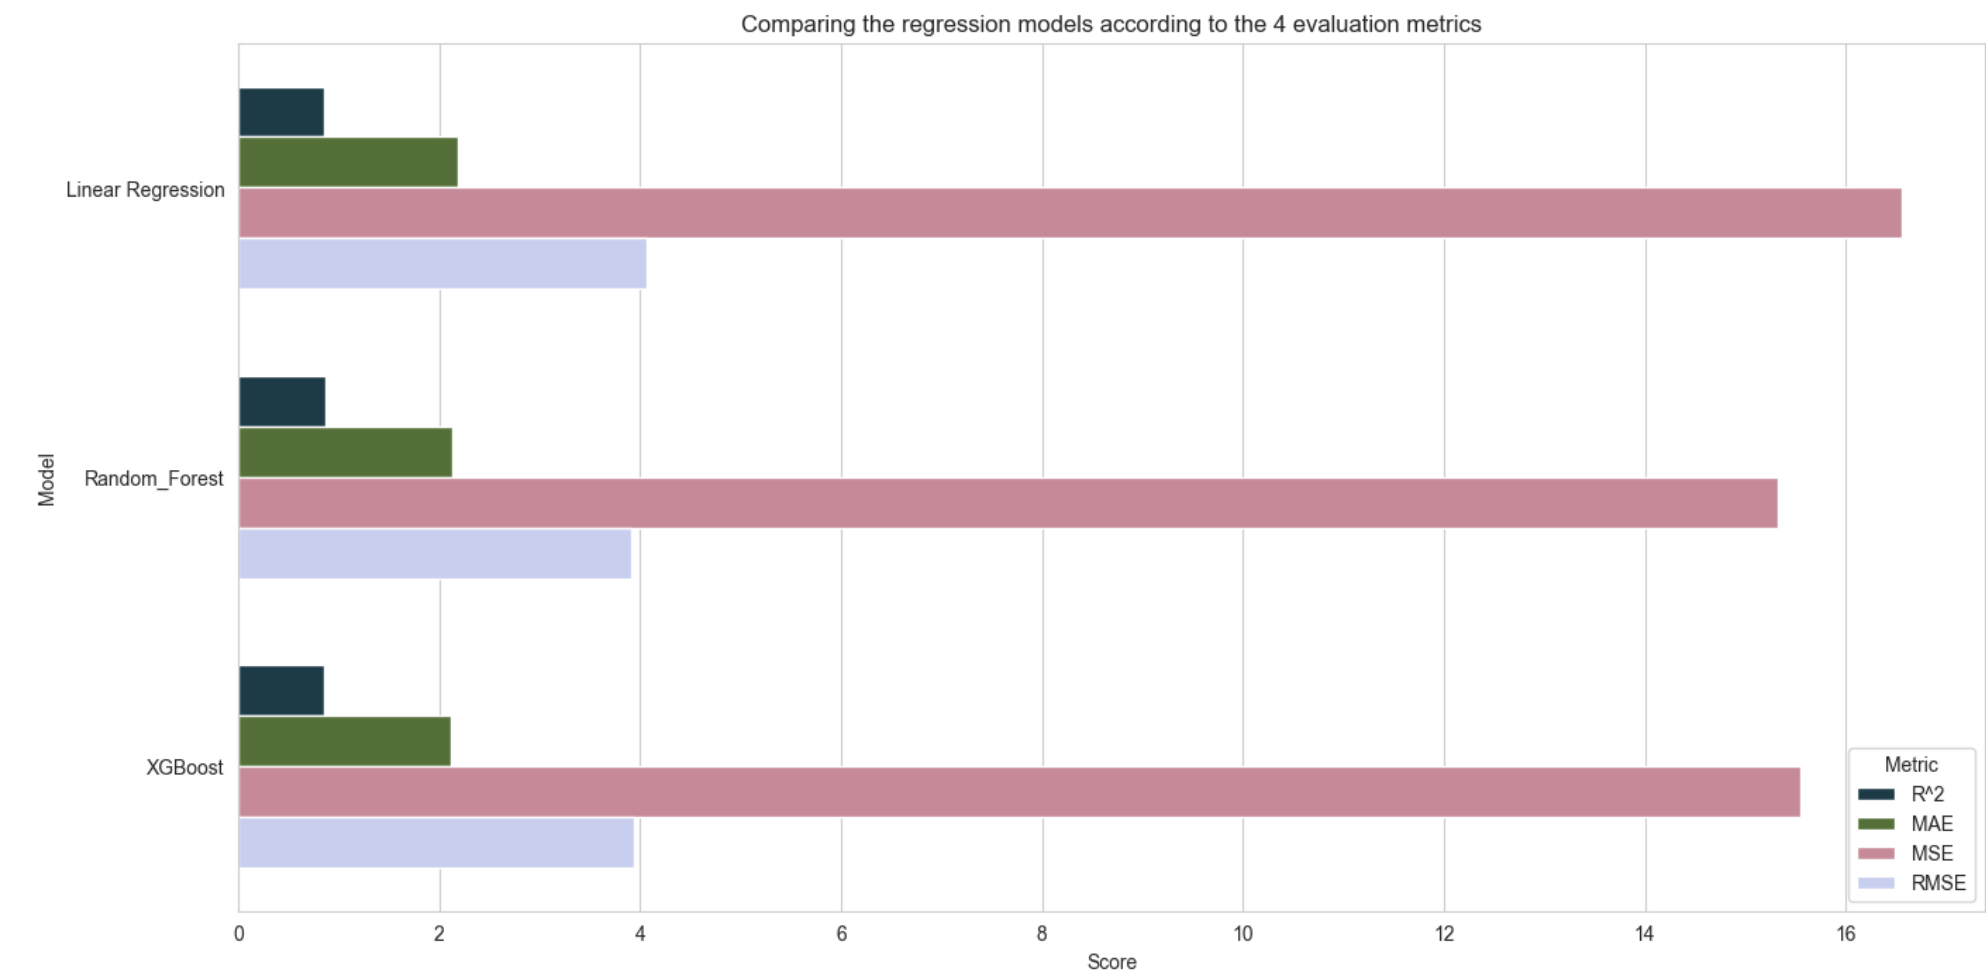
Comparing models

```
data = {"Model": ["Linear Regression", "Linear Regression", "Linear Regression", "Linear Regression",
"Random_Forest", "Random_Forest", "Random_Forest", "Random_Forest", "XGBoost", "XGBoost", "XGBoost",
"XGBoost"],\
"Metric": ["R^2", "MAE", "MSE", "RMSE", "R^2", "MAE", "MSE", "RMSE", "R^2", "MAE", "MSE", "RMSE"],\
"Score": [0.85, 2.19, 16.56, 4.07, 0.87, 2.13, 15.32, 3.91, 0.86, 2.12, 15.55, 3.94]}
```

```
comparing_models_ = pd.DataFrame(data)
comparing_models_
```

	Model	Metric	Score
0	Linear Regression	R^2	0.85
1	Linear Regression	MAE	2.19
2	Linear Regression	MSE	16.56
3	Linear Regression	RMSE	4.07
4	Random_Forest	R^2	0.87
5	Random_Forest	MAE	2.13
6	Random_Forest	MSE	15.32
7	Random_Forest	RMSE	3.91
8	XGBoost	R^2	0.86
9	XGBoost	MAE	2.12
10	XGBoost	MSE	15.55
11	XGBoost	RMSE	3.94

```
plt.figure( figsize= (16, 8))
fig = sns.barplot(data= comparing_models_, x= "Score", y= "Model", orient= "horizontal", hue= "Metric",
palette= "cubehelix", width= 0.7, dodge= True)
fig.set_title("Comparing the regression models according to the 4 evaluation metrics")
plt.show()
```



We could observe that the Random Forest Model performs the best based on the all metrics except for the mean absolute error.

Random forest model results

```
# Create a `results` dataframe
results = pd.DataFrame(data= {"actual": y_test["fare_amount"], "predicted": y_pred_test_ranfor.ravel()})
results.head(5)
```

	actual	predicted
5818	14.0	12.107306
18134	28.0	16.999474
4655	5.5	6.541329
7378	15.5	15.847178
13914	9.5	10.174875


```
results["residual"] = results["actual"] - results["predicted"]
results.head(5)
```

	actual	predicted	residual
5818	14.0	12.107306	1.892694
18134	28.0	16.999474	11.000526
4655	5.5	6.541329	-1.041329


```
7378      15.5  15.847178  -0.347178
13914      9.5  10.174875  -0.674875
```

Visualize random forest model results

```
# Create a scatterplot to visualize `predicted` over `actual`
plt.figure(figsize= (16, 8))
sns.scatterplot(data= results, x= "actual", y= "predicted", alpha= 0.7)
plt.plot([0, 60], [0, 60], c= "red", linewidth= "2")
plt.xticks(range(0, 86, 5))
plt.show()
```




```
# Visualize the distribution of the `residuals`
plt.figure(figsize= (16, 6))
sns.histplot(data= results, x= "residual", kde= True, bins= np.arange(-20, 20, 1))
plt.show()
```



```
# Calculate residual mean
print("Residual Mean =", round(results["residual"].mean(), 4))
```

Residual Mean = 0.0253

```
# Create a scatterplot of `residuals` over `predicted`
plt.figure(figsize= (16, 6))
sns.scatterplot(data= results, x= "predicted", y= "residual", alpha= 0.7)
plt.plot([0, 60], [0, 0], c= "red", linewidth= "2")
plt.show()
```



Feature importance

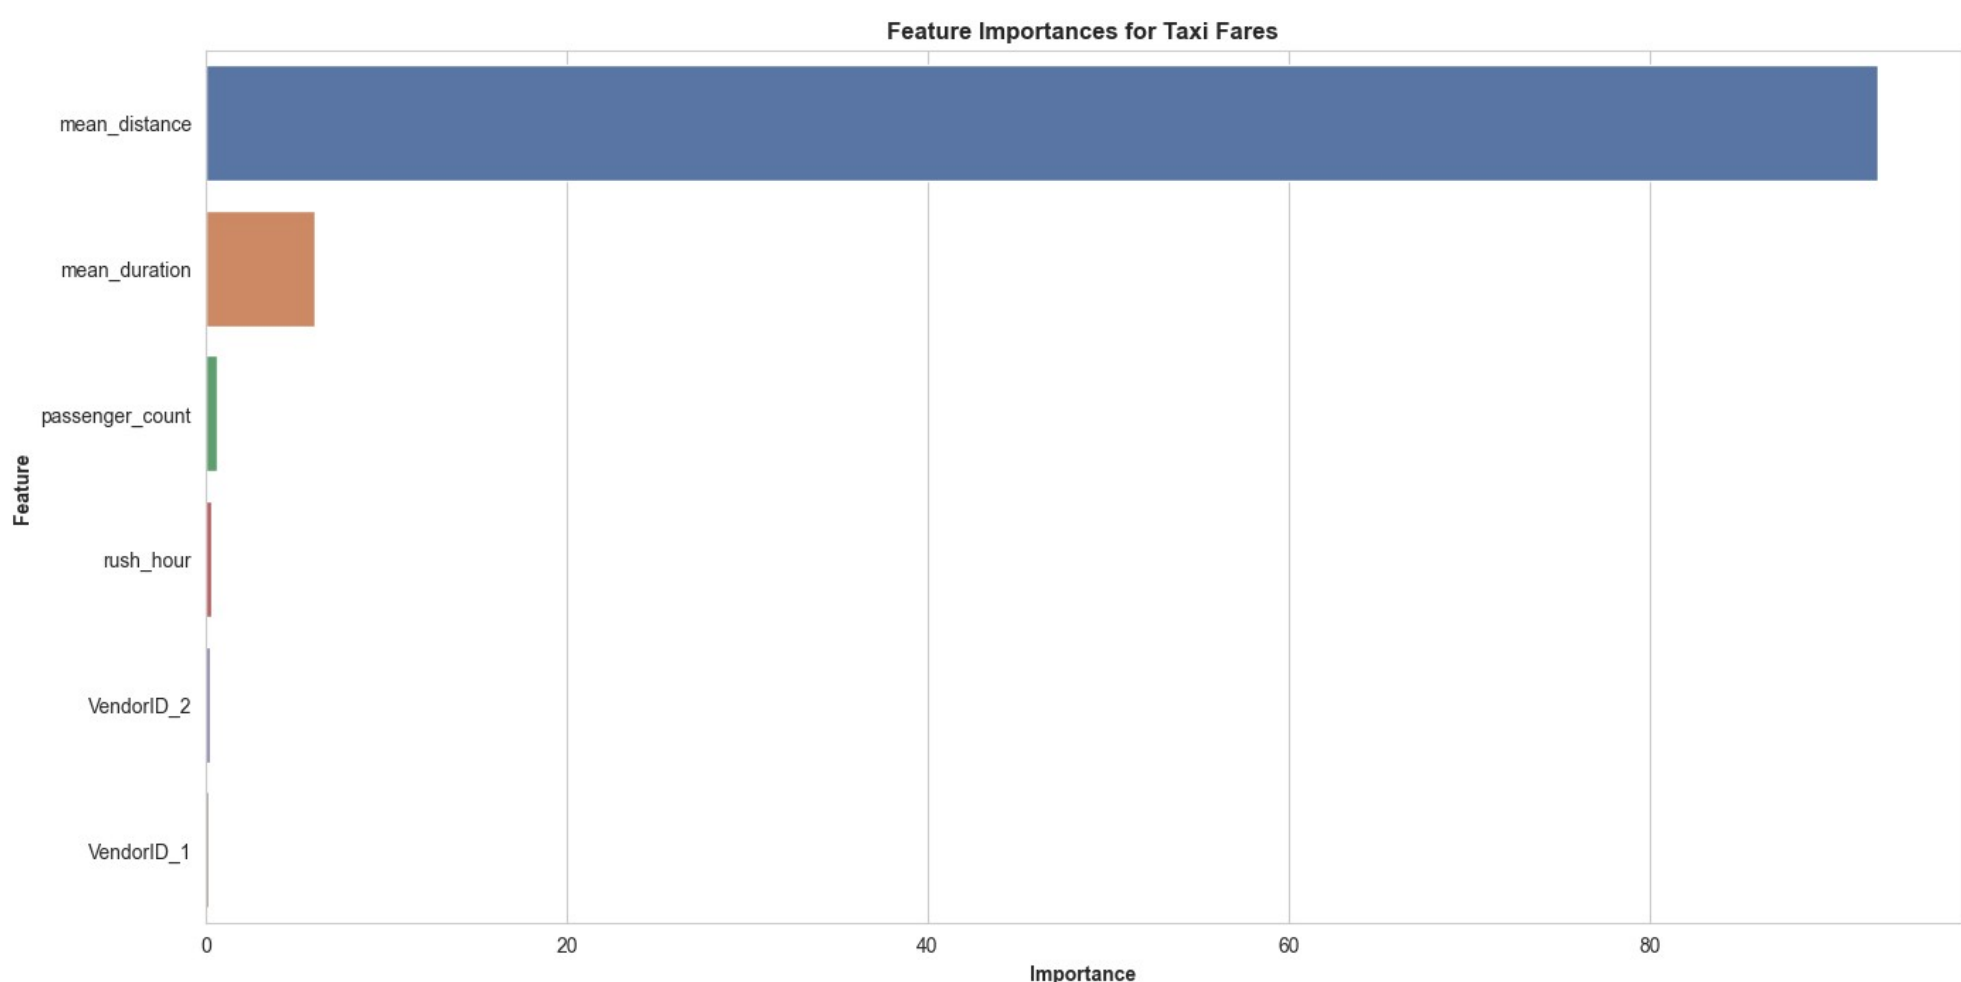
```
feature_importances = ranfor_reg.feature_importances_  
feature_importances
```

```
array([0.00325539, 0.00609513, 0.92578617, 0.06039367, 0.00206913,  
       0.00240051])
```

```
forest_importances = pd.DataFrame(list(X.columns), columns= ["feature"])  
forest_importances["importance"] = feature_importances * 100  
forest_importances = forest_importances.sort_values(by="importance", ascending= False)  
forest_importances
```

	feature	importance
2	mean_distance	92.578617
3	mean_duration	6.039367
1	passenger_count	0.609513
0	rush_hour	0.325539
5	VendorID_2	0.240051
4	VendorID_1	0.206913

```
plt.figure(figsize= (16, 8))  
sns.barplot(data= forest_importances, x= "importance", y= "feature", palette= "deep", orient= "h")  
plt.title("Feature Importances for Taxi Fares", fontweight= "bold")  
plt.ylabel("Feature", fontweight= "bold")  
plt.xlabel("Importance", fontweight= "bold")  
plt.show()
```



From the above graph, we observe that the mean distance of the trip is the most influential factor in estimating the taxi fares in advance with relative importance of 92.58%. In the future, adding more information on a rider's past behavior may also be beneficial in helping the stakeholder address their business problem.